

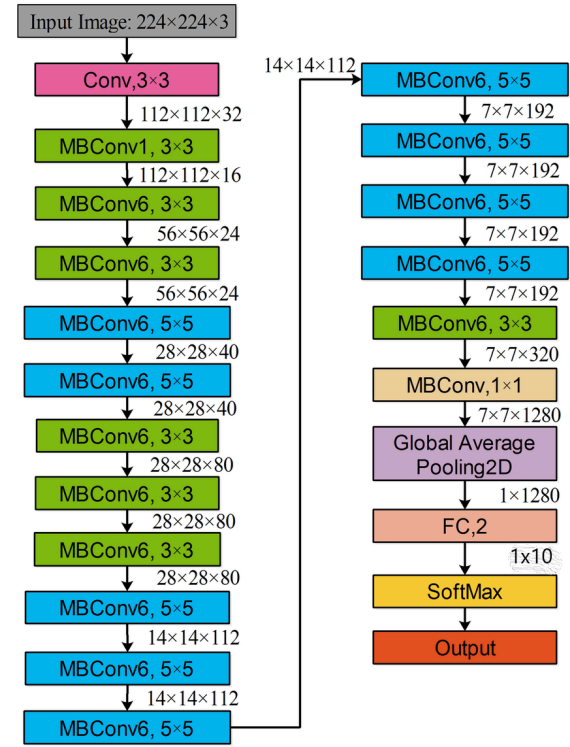
CIFAR-10 and STL-10 Image Classification

1. Introduction

Object recognition is a fundamental task in computer vision with applications in robotics and autonomous systems. This report compares deep learning and traditional computer vision approaches for image classification on two benchmark datasets: CIFAR-10 and STL-10. Both datasets have 10 object classes that largely overlap, with nine common objects: airplane, bird, cat, deer, dog, horse, ship, truck, and car (termed automobile in CIFAR-10), but differ in that CIFAR10 includes frog while STL-10 includes monkey. CIFAR-10 originally contained 50,000 training images and 10,000 test images [1], but was subsampled to 5,000 training images (500 per class) and 8,000 test images (800 per class) to match STL10's size and maintain balanced class representation. STL-10 is used in its entirety, with 5,000 training images (500 per class) and 8,000 test images (800 per class) [2], featuring higher resolution 96x96 images sourced from ImageNet. CIFAR-10 and STL-10 were chosen to evaluate the methods across different image characteristics: CIFAR-10s 32x32 images are low-resolution with centered subjects and minimal background clutter, while STL-10s images often include complex backgrounds and clutter, posing greater challenges for feature detection and classification. Using both datasets provides insights into how each approach handles varying levels of image complexity and resolution.

2. EfficientNet B0

The deep learning approach utilized EfficientNetB0, a CNN designed for efficiency through compound scaling of depth, width, and resolution [3]. Pre-trained on ImageNet, EfficientNetB0 leverages transfer learning, where pre-trained weights from a large dataset (ImageNet) are reused to improve performance on smaller target datasets (CIFAR-10, STL-10) by initializing with learned features. This architecture, with approximately 237 layers, processes inputs sequentially, as shown in **Figure 1**. Starting with a $224 \times 224 \times 3$ input, a 3×3 convolution layer extracts initial features. MBConv1 blocks at stages 2–3 apply depth wise separable convolutions for efficient feature extraction, followed by batch normalization to stabilize training and Swish activation for smooth nonlinearity. MBConv6 blocks at stages 4–6 deepen the network, using skip connections to enhance gradient flow. Stage 7 applies a 1×1 convolution, followed by global average pooling 2D to a 1×1280 vector, reducing overfitting. The architecture ends with two fully connected (FC) layers and softmax for ten object classes.



2.2 Hyperparameter Exploration

Hyperparameter tuning involved a grid search to systematically explore learning rates, dropout rates, and batch sizes. Learning rates (0.001, 0.0005, 0.0001) were chosen based on typical transfer learning ranges for EfficientNetB0, starting with 0.001 to ensure rapid convergence of the pretrained model while testing lower rates to refine optimization and prevent overshooting on small datasets. Dropout rates (0.2, 0.3, 0.4) addressed overfitting risks, with 0.2 as a conservative baseline to maintain model capacity, increasing to 0.4 to counter the high risk of overfitting. Batch sizes (32, 64, 128) balanced gradient stability and computational efficiency, with 32 aligning with typical transfer learning setups for small datasets, 64 as a middle ground, and 128 testing larger batches for improved generalization. The Adam optimizer, chosen for its adaptive learning rates, optimized convergence; clipnorm=1.0 capped gradient norms to prevent instability with small batches, and weight decay=1e-4 promoted generalization by penalizing large weights. Training ran for 25 epochs, which proved adequate as the majority of runs stopped early (typically within 10–15 epochs) due to early stopping (patience of 5), indicating efficient convergence.

The best hyperparameter configurations are then fine-tuned by adhering to the standard EfficientNetB0 transfer learning protocol. Initially, all layers except the custom head were frozen. The last 25 layers were unfrozen for CIFAR-10 to adapt to its simpler, centered images, while the last 10 layers were unfrozen for STL-10 to handle its complex backgrounds, balancing adaptation and stability. Batch normalization layers remained frozen to preserve pre-trained statistics. A ReduceLROnPlateau scheduler adjusted the learning rate from 0.001, with fine-tuning extending to 40 epochs and early stopping (patience of 5). Implemented in Python with TensorFlow and Keras, the approach used a fixed seed (42) and epoch timing callbacks, ensuring reproducibility and monitoring.

2.3 Hyperparameter Analysis

The baseline EfficientNetB0 (learning rate=0.001, dropout=0.2, batch size=32) was trained and achieved 78.97% training accuracy, 82.70% validation accuracy, and 79.70% test accuracy. Augmentation drove generalization, but upscaling 32×32 images to 224×224 introduced artifacts, impacting classes like “cat” and “dog”. STL-10 achieved 95.91% training, 94.60% validation, and 95.23% test accuracy. **Table 1** summarizes the hyperparameter tuning results, highlighting the best and worst configurations and performance ranges for both datasets.

Table 1: Key configurations from the 27-configuration grid search

Learning Rate	0.001	0.001	0.0005	0.0001
Dropout Rate	0.2	0.4	0.3	0.4
Batch Size	128	64	32	128
CIFAR-10 Test Accuracy (%)	82.27 (Best)	81.05	80.12	79.09 (Worst)
STL-10 Test Accuracy (%)	95.74 (Best)	95.3	95.2	95.06 (Worst)

For CIFAR-10, the best configuration achieved a test accuracy of 82.27%, improving the baseline test accuracy by 2.57%. The test accuracy spanned 79.09% to 82.27%, with a mean of approximately 80.98% and a standard deviation of 0.94%. The variance reflects CIFAR-10’s sensitivity to hyperparameters, driven by upscaling artifacts and class overlaps. For STL-10, the best configuration achieved a test accuracy of 95.74%, a 0.51% improvement over the baseline. Test accuracies ranged from 95.06% to 95.74%, with a mean of 95.45% and a standard deviation of 0.22%. The low variance indicates robust performance, attributed to STL-10’s higher resolution and ImageNet-derived features.

Hyperparameter analysis shows that a learning rate of 0.001 outperformed lower rates (e.g., 79.09% for CIFAR-10, 95.30% for STL-10 at 0.0001), as higher rates better navigate the initial loss landscape. A dropout rate of 0.2 optimized generalization, with higher rates (0.4) reducing accuracies to 81.05% for CIFAR-10 and 95.06% for STL-10 due to excessive regularization. A batch size of 128 enhanced gradient stability, particularly for CIFAR-10, where noisy gradients benefit from larger batches .

2.4 Fine-Tuning Performance

Fine-tuning tailored the optimal EfficientNetB0 configurations to each dataset’s unique characteristics, leveraging dataset-specific augmentation and partial unfreezing of layers to enhance performance. For CIFAR-10, the fine-tuned model achieved a test accuracy of 87.35%. This 9.65% improvement over the baseline reflects significant adaptation to CIFAR-10’s low-resolution challenges. For STL-10, fine-tuning achieved a test accuracy of 95.91%. The classification report shown in **Table 2.1** provides per-class insights of the fine-tuned model on CIFAR-10, where the macro-averaged precision, recall, and F1-score are 0.87, with “cat” having the lowest F1-score of 0.76 due to confusion with “dog” (F1=0.82). This is likely due to overlapping visual features like fur and shape. Stronger classes, such as “ship” (F1=0.94), benefit from distinct silhouettes.

The classification report shown in **Table 2.2** provides per-class insights of the fine-tuned model on STL-10, where the macro-averaged metrics are 0.96, with “bird” and “truck” excelling at F1=0.98, while “car” (F1=0.92, recall=0.89) underperforms slightly, likely due to variable backgrounds. CIFAR-10’s deeper fine-tuning and adaptive learning rate scheduling explain its larger accuracy gain compared to STL-10’s 0.68%, as more layers adapt to upscaling artifacts and class overlaps.

Table 2.1: Fine-tuned model classification report on CIFAR-10

Object	Precision	Recall	F1-Score
airplane	0.92	0.88	0.9
bird	0.91	0.95	0.93
automobile	0.85	0.84	0.85
cat	0.79	0.74	0.76
deer	0.87	0.81	0.84
dog	0.85	0.79	0.82
horse	0.81	0.95	0.87
frog	0.9	0.9	0.9
ship	0.93	0.95	0.94
truck	0.91	0.93	0.92
Macro Avg	0.87	0.87	0.87
Accuracy	0.87		

Table 2.2: Fine-tuned model classification report on STL-10

Object	Precision	Recall	F1-Score
airplane	0.97	0.97	0.97
bird	0.99	0.96	0.98
car	0.94	0.89	0.92
cat	0.94	0.97	0.95
deer	0.93	0.91	0.92
dog	0.93	0.96	0.95
horse	0.95	0.96	0.96
monkey	0.96	0.98	0.97
ship	0.95	0.96	0.96
truck	0.98	0.97	0.98
Macro avg	0.96	0.96	0.96
Accuracy	0.96		

The fine-tuning results provide granular insights into convergence as visualized in **Figure 2**. For CIFAR-10, training accuracy improved from 48.26% to 97.85%, with validation accuracy stabilizing at 87.90%. The ReduceLROnPlateau callback, reducing the learning rate four times due to stagnant validation loss, enabled precise weight updates and drove validation accuracy from 85.60% at Epoch 6 to 88.10%. Early stopping at Epoch 20 prevented overfitting. STL-10’s training accuracy rose from 31.74% to

98.04% by Epoch 23, with validation accuracy peaking at 96.10% at Epoch 19. The fixed learning rate of 0.0001 ensured stability, but the lack of scheduling limited further gains. Generalization was maintained by the Early stopping at Epoch 23.

CIFAR-10’s augmentation, incorporating 10% zoom, enhanced robustness to upscaling distortions by simulating variable object scales and contributed to the jump from 82.27% to 87.35% test accuracy. STL-10’s broader $\pm 20^\circ$ rotation supported generalization without compromising its richer features, yielding high F1-scores.

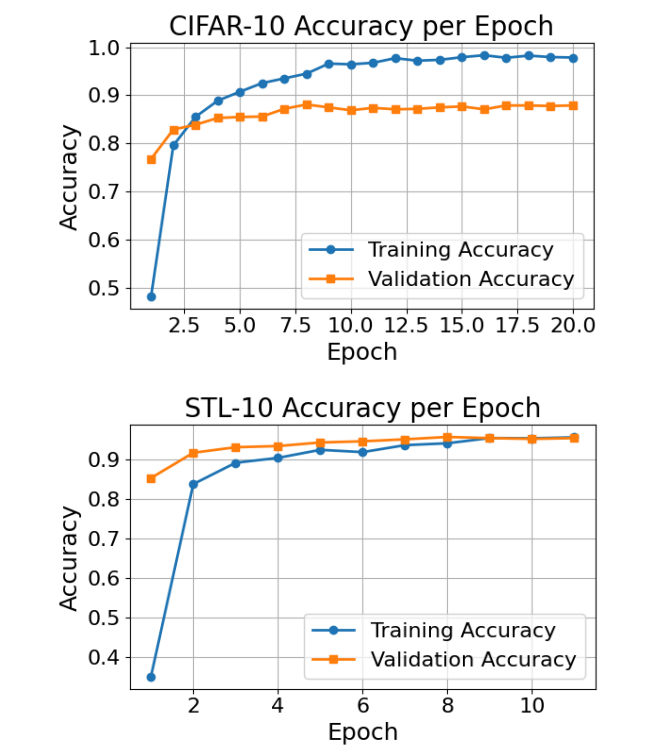


Fig. 2. Fine-Tuned Model performance on both datasets

2.5 Assessment and Future Directions

The fine-tuned results underscore EfficientNetB0’s adaptability to diverse datasets. STL-10’s 95.91% test accuracy and balanced classification metrics in **Table 2.2** highlight its suitability for transfer learning, leveraging 96×96 resolution and ImageNet-derived features. The low variance (standard deviation $\approx 0.2\%$), and rapid convergence (validation accuracy reaching 95.70% by Epoch 9 in **Figure 2**) indicate that the pretrained weights require minimal adjustment. The slight underperformance on “car” (F1=0.92) suggests background clutter may obscure foreground features, which could be addressed by integrating attention mechanisms such as CBAM to prioritize relevant image regions. Adding zoom augmentation, as used in CIFAR-10, could enhance STL-10’s robustness to scale variations, potentially pushing accuracy beyond 96%.

CIFAR-10’s fine-tuned 87.35% test accuracy, a 9.65% improvement over the baseline, reflects robust adaptation to low-resolution challenges, driven by unfreezing 25 layers, adaptive learning rate scheduling, and tailored augmentation. The classification report in **Table 2.1** highlights persistent confusion between “cat” ($F1=0.76$) and “dog” ($F1=0.82$), where visual similarities lead to misclassifications, while classes like “ship” ($F1=0.94$) benefit from distinct features. The higher variance (standard deviation $\approx 0.94\%$) and slower convergence (validation accuracy peaking at 88.10% by Epoch 8 as shown in **Figure 1**) underscore the dataset’s complexity, exacerbated by upscaling artifacts. The extended 25-epoch training, compared to fewer epochs in earlier experiments, allowed sustained validation accuracy gains and explains the improved 87.35%.

Future work could enhance CIFAR-10’s performance by applying texture-enhancing augmentations such as edge detection filters to emphasize distinguishing features like fur patterns, potentially improving “cat” and “dog” F1-scores. For STL-10, ensemble methods combining EfficientNetB0 with ResNet50 could boost “car” performance by leveraging complementary feature representations. Exploring alternative architectures, such as ResNet50 alone, could validate EfficientNetB0’s suitability or reveal better alternatives. Testing higher learning rates (e.g., 0.002) or larger batch sizes (e.g., 256) in future hyperparameter searches could uncover configurations better suited to CIFAR-10’s complex feature space or STL-10’s stable performance.

3. Traditional CV

This section covers two traditional computer vision approaches for object classification, employing the Bag of Words framework and comparing SIFT with Difference of Gaussians (DoG) against SIFT with Harris-Laplace. The BoW pipeline in both approaches comprise of five stages: (1) sparse keypoint detection to identify salient regions, (2) descriptor extraction to characterize these keypoints, (3) codebook generation to cluster descriptors, (4) histogram representation to encode image features, and (5) classification to assign category labels.

Sparse sampling is particularly effective for STL-10’s complex, cluttered scenes and CIFAR-10’s compact yet distinct object patterns, as it focuses on visually significant areas rather than dense, uniform sampling. Descriptors extracted from keypoints are processed using PCA to reduce dimensionality while retaining 90% of the variance. PCA enhances computational efficiency and ensures that the most informative components of the descriptors are retained, mitigating redundancy and noise. Codebooks are constructed using MiniBatchKMeans with a batch size of 4096, selected over traditional K-means for its scalability and reduced memory footprint when clustering vast

descriptor sets from both datasets. The batch size of 4096 was chosen as a practical compromise to ensure stable clustering without exhausting computational resources. Vocabulary sizes of 750, 1000, and 1250 were explored via grid search, allowing sufficient granularity to capture STL-10’s diverse object appearances while promoting generalization across CIFAR-10’s simpler, more uniform categories.

To address BoW’s limitation in capturing spatial relationships, Spatial Pyramid Matching at level 1 is incorporated, providing a coarse representation of spatial layout without inflating dimensionality, which is a key advantage for CIFAR-10’s low resolution and STL-10’s cluttered images. L2 normalization is applied to the resulting histograms to ensure scale invariance and enhance the robustness of the subsequent classification step by mitigating the impact of varying feature magnitudes. For classification, SVM with a RBF kernel is employed due to its practical efficacy in high-dimensional, non-linear feature spaces produced by BoW histograms.

The grid search was performed manually on the test set by systematically evaluating all combinations of hyperparameters. For SIFT with DoG, the grid included contrast thresholds (0.02, 0.04, 0.08) to control keypoint density, edge thresholds (7.5, 10, 12.5) to filter edge-related noise, vocabulary sizes (750, 1000, 1250) to adjust codebook granularity, and SVM C values (0.1, 1.0, 10.0) to tune regularization. For SIFT with Harris-Laplace, the grid was extended to include additional parameters: block sizes (3 and 4 for STL-10; 1, 2, and 3 for CIFAR-10) to adjust the scale of corner detection, sensitivity k values (0.05 and 0.07 for STL-10; 0.04 and 0.06 for CIFAR-10) to control corner strength, and sigma values ([1.0, 2.0, 4.0] for STL-10; [1.0, 1.5, 2.0] for CIFAR-10) to manage scale-space smoothing. The grid also maintained the same contrast thresholds (0.02, 0.04, 0.08), edge thresholds (7.5, 10, 12.5), vocabulary sizes (750, 1000, 1250), and SVM C values (0.1, 1.0, 10.0) as used in the DoG-based configuration.

3.1 DOG & Harris-Laplace

SIFT with DoG identifies keypoints by detecting extrema in the DoG scale space, an efficient approximation of the Laplacian-of-Gaussian (LoG). This method ensures scale invariance, making it well-suited for STL-10’s diverse object scales and CIFAR-10’s smaller, uniform objects. DoG operates by subtracting Gaussian-blurred images at adjacent scales, highlighting regions with significant contrast or texture changes, and is ideal for sparse, robust keypoint detection. The grid search tuned contrast thresholds to adjust keypoint density, with lower values retaining more points in STL-10’s low-contrast scenes and higher values prioritizing quality in CIFAR-10’s simpler images. Edge thresholds filter noisy keypoints near edges, with higher values reducing clutter in STL-10 and lower values preserving detail in

CIFAR-10. Descriptors, computed over 16x16 regions with orientation normalization provide rotation-invariant gradient patterns that are critical for both datasets’ viewpoint variations.

SIFT with Harris-Laplace integrates multi scale Harris corner detection with LoG scale selection, combining precise spatial localization with scale invariance. Harris identifies corners via a second-moment matrix, detecting regions with multidirectional intensity changes, filtered at a 1% response threshold for sparsity. LoG assigns scales by locating extrema across smoothed scales and is proficient for handling STL-10’s larger objects and CIFAR-10’s finer details. The previously mentioned grid search tailored both CV approaches to each dataset’s scale and complexity.

3.2 SIFT (DoG) Results

The best SIFT with DoG configuration on CIFAR-10 achieved an accuracy of 0.2559 with a contrast threshold of 0.04, edge threshold of 10.0, vocabulary size of 750, and SVM parameters C=1.0. **Table 4.1** summarizes the accuracy for a subset of configurations, selected to include the best performance and illustrate hyperparameter trends.

Table 4.1: Accuracy of SIFT with DoG on CIFAR-10 for a Subset of Configurations

Contrast Thresh.	Edge Thresh.	Vocab Size	C=0.1	C=1.0	C=10.0
0.02	7.5	750.0	0.2114	0.2519	0.2456
0.02	10.0	1000.0	0.2124	0.2547	0.2509
0.04	10.0	750.0	0.2221	0.2559	0.251
0.04	12.5	1250.0	0.2171	0.2542	0.2449
0.08	7.5	1000.0	0.2161	0.2442	0.233

The classification report for the best SIFT with DOG configuration in **Table 4.2** shows modest performance, with airplane (F1=0.34) and truck (F1=0.32) performing relatively well, while bird and deer (F1=0.18) are the most challenging.

Table 4.2: Classification Report for Best SIFT with DoG Configuration on CIFAR-10

Class	Precision	Recall	F1-Score
Airplane	0.37	0.32	0.34
Automobile	0.34	0.28	0.31
Bird	0.20	0.17	0.18
Cat	0.19	0.20	0.20
Deer	0.16	0.20	0.18
Dog	0.23	0.26	0.25
Frog	0.21	0.21	0.21
Horse	0.28	0.26	0.27
Ship	0.33	0.32	0.33
Truck	0.29	0.34	0.32
Macro Avg	0.26	0.26	0.26
Weighted Avg	0.26	0.26	0.26

The best SIFT with DoG configuration on STL-10 achieved an accuracy of 0.3877 with a contrast threshold of 0.02, edge threshold of 12.5, vocabulary size of 750, and SVM parameters C=1.0. **Table 4.3** summarizes the accuracy for a subset of configurations, selected to include the best performance and illustrate hyperparameter trends for brevity.

Table 4.3: Accuracy of SIFT with DoG on STL-10 for a Subset of Configurations

Contrast Thresh.	Edge Thresh.	Vocab Size	C=0.1	C=1.0	C=10.0
0.02	7.5	1000.0	0.2731	0.3695	0.3678
0.02	12.5	750.0	0.2825	0.3877	0.3812
0.04	10.0	750.0	0.2864	0.3759	0.3685
0.04	12.5	1050.0	0.2484	0.3777	0.3723
0.08	10.0	1250.0	0.2301	0.3438	0.3364

The classification report for the best configuration in **Table 4.4** shows improved performance compared to CIFAR10, with airplane (F1=0.56) and ship (F1=0.49) performing strongly, while automobile (F1=0.24) and dog (F1=0.28) are less accurate.

Table 4.4: Classification Report for Best SIFT with DoG Configuration on STL-10

Class	Precision	Recall	F1-Score
Airplane	0.58	0.55	0.56
Automobile	0.30	0.20	0.24
Bird	0.42	0.49	0.46
Cat	0.35	0.33	0.34
Deer	0.34	0.41	0.37
Dog	0.27	0.29	0.28
Frog	0.35	0.35	0.35
Horse	0.30	0.32	0.31
Ship	0.49	0.50	0.49
Truck	0.49	0.44	0.46
Macro Avg	0.39	0.39	0.39
Weighted Avg	0.39	0.39	0.39

4.3 SIFT (Harris-Laplace) Results

The best SIFT with Harris-Laplace configuration on CIFAR-10 achieved an accuracy of 0.4031 with a block size of 3, k=0.06, contrast threshold of 0.02, edge threshold of 7.5, vocabulary size of 1250, and SVM parameters C=1.0. **Table 4.5** summarizes the accuracy for a subset of configurations, selected to include the best performance and illustrate hyperparameter trends for brevity.

Table 4.5: Classification Report for Best SIFT with DoG Configuration on CIFAR-10

Block Size	k	Contrast	Edge	Vocab Size	C=0.1	C=1.0
2	0.04	0.02	7.5	1250	0.3084	0.3915
2	0.06	0.02	7.5	1250	0.2990	0.3760
3	0.04	0.08	12.5	1000	0.3142	0.4006
3	0.06	0.02	7.5	1250	0.3114	0.4031
3	0.06	0.04	10.0	1250	0.3114	0.4031

The classification report for the best configuration in **Table 4.6** shows balanced performance, with ships (F1=0.53) and automobiles (F1=0.49) performing well, and cats (F1=0.28) and deer (F1=0.29) showing lower performance.

Table 4.6: Classification Report for Best SIFT with Harris-Laplace Configuration on CIFAR-10

Class	Precision	Recall	F1-Score
Airplane	0.49	0.42	0.45
Automobile	0.50	0.47	0.49
Bird	0.33	0.29	0.31
Cat	0.27	0.29	0.28
Deer	0.31	0.27	0.29
Dog	0.33	0.39	0.36
Frog	0.37	0.45	0.41
Horse	0.51	0.38	0.43
Ship	0.51	0.55	0.53
Truck	0.45	0.51	0.48
Macro Avg	0.41	0.40	0.40
Weighted Avg	0.41	0.40	0.40

The best SIFT with Harris-Laplace configuration on STL-10 achieved an accuracy of 0.5051 with a block size of 4, k=0.05, contrast threshold of 0.04, edge threshold of 7.5, vocabulary size of 1000, and SVM parameters C=10.0. **Table 4.7** summarizes the accuracy for a subset of configurations, selected to include the best performance and illustrate hyperparameter trends for brevity.

Table 4.7: Accuracy of SIFT with Harris-Laplace on STL-10 for a Subset of Configurations

k	Contrast	Edge	Vocab Size	C=0.1	C=1.0	C=10.0
0.05	0.02	7.5	1000	0.4015	0.4818	0.4843
0.05	0.04	10.0	1250	0.4041	0.4799	0.4816
0.07	0.02	7.5	1250	0.3970	0.4745	0.4714
0.05	0.02	10.0	1000	0.4111	0.5021	0.5051
0.05	0.04	7.5	1000	0.4111	0.5021	0.5051

The classification report for the best configuration shown in **Table 4.8** shows strong performance for ship (F1=0.67) and airplane (F1=0.63), driven by their distinct edges, while cat (F1=0.35) and dog (F1=0.36) are less accurate due to texture complexity. The balanced macro average (F1=0.50) and accuracy (0.5051) reflect robust performance, enhanced by the larger block size (4) and moderate vocabulary size (1000) on STL-10s 96×96 images.

Table 4.8: Classification Report for Best SIFT with Harris-Laplace Configuration on STL-10

Class	Precision	Recall	F1-Score
Airplane	0.59	0.69	0.63
Automobile	0.42	0.43	0.42
Bird	0.59	0.60	0.59
Cat	0.36	0.33	0.35
Deer	0.46	0.48	0.47
Dog	0.37	0.35	0.36
Frog	0.55	0.55	0.55
Horse	0.41	0.37	0.39
Ship	0.66	0.69	0.67
Truck	0.61	0.57	0.59
Macro Avg	0.50	0.51	0.50
Weighted Avg	0.50	0.51	0.50

For CIFAR-10, SIFT with Harris-Laplace (accuracy: 0.4031) significantly outperforms SIFT with DoG (accuracy: 0.2559). Harris-Laplace's multi-scale corner detection captures more discriminative keypoints in the low resolution 32×32 images. The optimal DoG configuration uses a moderate contrast threshold (0.04) and edge threshold (10.0) with a smaller vocabulary size (750), while Harris-Laplace benefits from a larger vocabulary size (1250), indicating a richer visual dictionary enhances its discriminative power. For STL-10, SIFT with Harris-Laplace (accuracy: 0.5051) outperforms SIFT with DoG (accuracy: 0.3877). The larger block size (4) and moderate k (0.05) in the optimal Harris-Laplace configuration likely improve keypoint stability on STL-10s higher-resolution 96×96 images. The vocabulary size of 1000 strikes a balance between feature granularity and computational efficiency, with C=10.0 suggesting a tighter margin for SVM classification. The DoG configuration, with a lower contrast threshold (0.02) and smaller vocabulary size (750), is less effective, possibly due to fewer robust keypoints.

4.3 Comparison and Interpretation

Across both datasets, Harris-Laplace consistently outperforms DoG, likely due to its ability to detect corner like features across multiple scales, which is advantageous for both low- and high-resolution images. STL-10s higher resolution benefits both methods, but Harris-Laplace's performance gap is more pronounced, suggesting better adaptability to larger images. The SVM parameter C=1.0 or C=10.0 yields the highest accuracies, with C=0.1 underfitting. Classification reports show that classes with distinct edges (e.g., airplane, ship) achieve higher F1-scores, while textured classes (e.g., cat, deer) are challenging, particularly on CIFAR-10. The STL-10 Harris Laplace classification reports confirm this trend, with ship and airplane excelling, while cat and dog remain difficult.

5. Deep learning V.S. Traditional CV

The object recognition results demonstrate that the deep learning approach using EfficientNet B0 significantly outperforms traditional computer vision methods employing SIFT with Difference of Gaussians (DoG) and Harris-Laplace on both CIFAR-10 and STL-10 datasets. EfficientNet B0 achieved test accuracies of 87.35% on CIFAR-10 and 95.91% on STL-10, compared to 40.31% (Harris-Laplace) and 25.59% (DoG) on CIFAR-10, and 50.51% (Harris-Laplace) and 38.77% (DoG) on STL-10. This performance gap is driven by EfficientNet B0's use of transfer learning, leveraging pre-trained ImageNet weights to extract rich, adaptable features, while SIFT methods rely on hand-crafted, gradient-based features that struggle with CIFAR-10's low-resolution 32x32 images (F1-scores of 0.26 for DoG, 0.40 for Harris-Laplace) and only moderately improve on STL-10's 96x96 images (F1-scores of 0.39 for DoG, 0.50 for Harris-Laplace). Deep learning excels across all classes, capturing both edge and texture patterns (e.g., F1=0.97 for airplane, 0.95 for cat on STL-10), whereas traditional methods favor edge-distinct classes like ship (F1=0.67, Harris-Laplace on STL-10) but falter on textured classes like cat (F1=0.35) and dog (F1=0.36). Data augmentation and end-to-end optimization further enhance EfficientNet B0's generalization, unlike the manual tuning required for SIFT's feature extraction and SVM classification. While traditional CV offers computational efficiency, deep learning's superior adaptability and feature learning make it far more effective for complex image classification tasks in robotics and autonomous systems.

6. State of the Art in Computer Vision for Robotics

Contemporary deep learning methods in robotic vision encompass a diverse range of architectures tailored to specific tasks. CNNs remain the backbone of many vision-based robotic systems, excelling in object detection and image classification. Pre-trained models such as ResNet and DenseNet are commonly fine-tuned for robotic applications, achieving robust performance across various environments [4]. More recently, transformer architectures have been adapted for vision tasks through models like Vision Transformers (ViTs). These models process images as sequences of patches, offering scalability and robustness for complex scene understanding in robotics [5]. Additionally, deep generative models, including Variational Autoencoders, Generative Adversarial Networks, and Diffusion Models, are gaining traction for generating synthetic data and learning complex distributions for tasks like robotic grasping. These deep learning methods support a wide array of robotic applications. Object detection and recognition are vital for identifying and interacting with objects in real-time, with state-of-the-art models like YOLO and Faster R-CNN achieving high accuracy in diverse

settings [6]. In navigation and path planning, deep learning enhances visual Simultaneous Localization and Mapping (SLAM) and obstacle avoidance, enabling robots to operate in dynamic environments [7].

Despite the transformative impact of deep learning, several challenges persist. One significant limitation is the requirement for large, annotated datasets, which are costly and time-consuming to acquire in robotics due to the variability of environments and tasks. To address this, researchers are exploring transfer learning, where models pre-trained on large datasets are fine-tuned on smaller, task-specific datasets, and synthetic data generation using simulation environments [8]. Generalization across different environments remains a concern, as models trained in specific settings may not perform well in new scenarios due to variations in lighting, object appearances, or other factors. Domain adaptation and robust training methods are active areas of research to improve model generalization [9]. Recent advancements and emerging trends are shaping the future of robotic vision. Multimodal learning, which integrates vision with other sensory inputs such as language or tactile feedback, is gaining traction to enhance robot understanding and interaction. Neural Radiance Fields (NeRF) are being explored for high-fidelity 3D scene reconstruction, which can provide detailed environmental models for robotic navigation and manipulation [10]. These trends are expected to drive significant progress in the field, making robots more autonomous and capable in diverse, real-world settings.

References

- [1] A. Krizhevsky, "Learning Multiple Layers of Features from Tiny Images," Jan. 2009, [Online]. Available: <https://www.cs.toronto.edu/~kriz/learning-features-2009-TR.pdf>
- [2] A. Coates, A. Ng, and H. Lee, "An Analysis of Single-Layer Networks in Unsupervised Feature Learning", in Proceedings of the Fourteenth International Conference on Artificial Intelligence and Statistics, 11–13 Apr 2011, vol. 15, pp. 215–223.
- [3] M. Tan and Q. V. Le, "EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks," arXiv.org, 2019, <https://arxiv.org/abs/1905.11946>
- [4] S. Sultana, Muhammad Mansoor Alam, Mazliham Mohd Su'ud, Jawahir Che Mustapha, and M. Prasad, "A Deep Dive into Robot Vision - An Integrative Systematic Literature Review Methodologies and Research Endeavor Practices," *ACM computing surveys*, vol. 56, no. 9, pp. 1–33, Apr. 2024, doi: <https://doi.org/10.1145/3648357>.
- [5] A. Dosovitskiy *et al.*, "An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale," *arXiv:2010.11929 [cs]*, Oct. 2020, Available: <https://arxiv.org/abs/2010.11929>
- [6] J. Redmon, S. Divvala, R. Girshick, and A. Farhadi, "You Only Look Once: Unified, Real-Time Object Detection," *arXiv.org*, Jun. 08, 2015, <https://arxiv.org/abs/1506.02640>
- [7] C. Cadena *et al.*, "Past, Present, and Future of Simultaneous Localization And Mapping: Towards the Robust-Perception Age," *IEEE Transactions on Robotics*, vol. 32, no. 6, pp. 1309–1332, Dec. 2016, doi: <https://doi.org/10.1109/TRO.2016.2624754>
- [8] J. Tobin, R. Fong, A. Ray, J. Schneider, W. Zaremba, and P. Abbeel, "Domain Randomization for Transferring Deep Neural Networks from Simulation to the Real World," arXiv:1703.06907 [cs], Mar. 2017, Available: <https://arxiv.org/abs/1703.06907>
- [9] E. Tzeng, J. Hoffman, K. Saenko, and T. Darrell, "Adversarial Discriminative Domain Adaptation," *openaccess.thecvf.com*, 2017, https://openaccess.thecvf.com/content_cvpr_2017/html/Tzeng_Adversarial_Discriminative_Domain_CVPR_2017_paper.html
- [10] B. Mildenhall, P. P. Srinivasan, M. Tancik, J. T. Barron, R. Ramamoorthi, and R. Ng, "NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis," arXiv:2003.08934 [cs], Aug. 2020, Available: <https://arxiv.org/abs/2003.08934>

Appendix A

1. CIFAR-10 CNN Experimentation

```
import os
import warnings
import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.preprocessing.image import
ImageDataGenerator
from tensorflow.keras.applications import
EfficientNetB0
from tensorflow.keras.applications.efficientnet
import preprocess_input
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import Callback
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import itertools
from sklearn.metrics import classification_report,
confusion_matrix
from sklearn.model_selection import
train_test_split
import time

# Suppress TensorFlow warnings
os.environ["TF_CPP_MIN_LOG_LEVEL"] = "3"
os.environ["TF_XLA_FLAGS"] =
"--tf_xla_auto_jit=-1"
tf.config.optimizer.set_jit(False)
warnings.filterwarnings("ignore",
category=UserWarning, module="keras")

# Set random seed for reproducibility
tf.random.set_seed(42)
np.random.seed(42)

# Define class names
cifar10_classes = ['airplane', 'automobile', 'bird',
'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
NUM_CLASSES = 10

# Custom callback for epoch timing
class EpochTimeCallback(Callback):
    def on_epoch_begin(self, epoch, logs=None):
```

```
        self.start_time = time.time()

    def on_epoch_end(self, epoch, logs=None):
        end_time = time.time()
        duration = end_time - self.start_time
        print(f"Epoch {epoch + 1} took {duration:.2f}
seconds")

# Load and preprocess CIFAR-10 with
subsampling training and test sets
def load_cifar10_data():
    """
    Loads CIFAR-10, subsamples to 5,000 train
    (500/class) and 8,000 test (800/class) (Kornblith
    et al., 2019).
    """
    try:
        (x_train, y_train), (x_test, y_test) =
cifar10.load_data()

        # Subsample training set to 5,000 (500 per
class)
        train_samples_per_class = 500
        x_train_sub, y_train_sub = [], []
        for cls in range(NUM_CLASSES):
            cls_indices = np.where(y_train[:, 0] ==
cls)[0]
            selected_indices =
np.random.choice(cls_indices,
train_samples_per_class, replace=False)

            x_train_sub.append(x_train[selected_indices])

        y_train_sub.append(y_train[selected_indices])
        x_train = np.vstack(x_train_sub)
        y_train = np.vstack(y_train_sub)

        # Subsample test set to 8,000 (800 per
class)
        test_samples_per_class = 800
        x_test_sub, y_test_sub = [], []
        for cls in range(NUM_CLASSES):
            cls_indices = np.where(y_test[:, 0] ==
cls)[0]
            selected_indices =
np.random.choice(cls_indices,
test_samples_per_class, replace=False)

            x_test_sub.append(x_test[selected_indices])

        y_test_sub.append(y_test[selected_indices])
        x_test = np.vstack(x_test_sub)
```



```

y_test = np.vstack(y_test_sub)

# Preprocess
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train = preprocess_input(x_train)
x_test = preprocess_input(x_test)
y_train = tf.keras.utils.to_categorical(y_train,
NUM_CLASSES)
y_test = tf.keras.utils.to_categorical(y_test,
NUM_CLASSES)

# Verify class distribution in test set
class_counts =
np.bincount(np.argmax(y_test, axis=1))
print("CIFAR-10 test set class distribution:",
class_counts)
assert all(count == test_samples_per_class
for count in class_counts), "Uneven class
distribution in CIFAR-10 test set"

print("CIFAR-10 x_train shape:",
x_train.shape, "y_train shape:", y_train.shape)
print("CIFAR-10 x_test shape:",
x_test.shape, "y_test shape:", y_test.shape)
return x_train, y_train, x_test, y_test
except Exception as e:
print(f"Error loading CIFAR-10: {str(e)}")
return None, None, None, None

# Build EfficientNetB0 model
def build_model(learning_rate=0.001,
dropout_rate=0.2):
"""
Constructs EfficientNetB0 with
BatchNormalization (Tan & Le, 2019; Ioffe &
Szegedy, 2015).
"""
base_model =
EfficientNetB0(weights='imagenet',
include_top=False, input_shape=(224, 224, 3))
base_model.trainable = False

model = models.Sequential([
layers.Input(shape=(32, 32, 3)),
layers.Resizing(224, 224),
base_model,
layers.GlobalAveragePooling2D(),
layers.Dense(128, activation='relu',
kernel_initializer='he_normal'),
layers.BatchNormalization(),
layers.Dropout(dropout_rate),

```

```

layers.Dense(NUM_CLASSES,
activation='softmax')
])

model.compile(optimizer=Adam(learning_rate=lea
rning_rate, clipnorm=1.0, weight_decay=1e-4),
loss='categorical_crossentropy',
metrics=['accuracy'])
return model

# Data augmentation
def create_data_generator():
"""
Applies augmentation with contrast (Shorten &
Khoshgoftaar, 2019).
Rotation  $\pm 20^\circ$ , shifts 10%, flips, brightness [0.8,
1.2].
"""
return ImageDataGenerator(
rotation_range=20,
width_shift_range=0.1,
height_shift_range=0.1,
horizontal_flip=True,
brightness_range=[0.8, 1.2],
fill_mode='nearest'
)

# Train and evaluate model
def train_model(model, x_train, y_train, x_val,
y_val, batch_size=32, epochs=25):
"""
Trains with augmentation, early stopping, and
epoch timing (Goodfellow et al., 2016).
"""
datagen = create_data_generator()
datagen.fit(x_train)

early_stopping =
tf.keras.callbacks.EarlyStopping(
monitor='val_loss', patience=5,
restore_best_weights=True
)
epoch_timer = EpochTimeCallback()

history = model.fit(
datagen.flow(x_train, y_train,
batch_size=batch_size),
validation_data=(x_val, y_val),
epochs=epochs,
callbacks=[early_stopping, epoch_timer],
verbose=1

```

```

)
return history

# Plot training results
def plot_results(history, title='CIFAR-10
EfficientNetB0 Performance'):
    """
    Generates loss/accuracy plots for report
    visualization (Marking Criteria 2).
    """
    plt.figure(figsize=(12, 4))
    plt.subplot(1, 2, 1)
    plt.plot(history.history['loss'], label='Train Loss')
    plt.plot(history.history['val_loss'],
label='Validation Loss')
    plt.title(f'{title} - Loss')
    plt.xlabel('Epoch')
    plt.ylabel('Loss')
    plt.legend()

    plt.subplot(1, 2, 2)
    plt.plot(history.history['accuracy'], label='Train
Accuracy')
    plt.plot(history.history['val_accuracy'],
label='Validation Accuracy')
    plt.title(f'{title} - Accuracy')
    plt.xlabel('Epoch')
    plt.ylabel('Accuracy')
    plt.legend()

    plt.savefig(f'{title.lower().replace(" ", "_").png}')
    plt.close()

# Generate confusion matrix
def save_confusion_matrix(y_true, y_pred,
classes, filename):
    """
    Saves confusion matrix for detailed class
    analysis (Marking Criteria 2).
    """
    cm = confusion_matrix(y_true, y_pred)
    plt.figure(figsize=(10, 8))
    plt.imshow(cm, interpolation='nearest',
cmap=plt.cm.Blues)
    plt.title('Confusion Matrix')
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)
    plt.xlabel('Predicted Label')
    plt.ylabel('True Label')
    plt.tight_layout()

```

```

plt.savefig(filename)
plt.close()

# Run baseline and grid search for CIFAR-10
def run_cnn_grid_search():
    """
    Runs baseline and grid search for CIFAR-10
    with BatchNormalization (Sokolova & Lapalme,
    2009).
    Baseline: LR=0.001, Dropout=0.2, BS=32.
    Grid: LR=[0.001, 0.0005, 0.0001],
    Dropout=[0.2, 0.3, 0.4], BS=[32, 64, 128].
    """
    x_train, y_train, x_test, y_test =
load_cifar10_data()
    classes = cifar10_classes

    if x_train is None:
        return

    x_train, x_val, y_train, y_val = train_test_split(
        x_train, y_train, test_size=0.2,
        random_state=42, stratify=y_train
    )
    print("CIFAR-10 Train shape:", x_train.shape,
"Validation shape:", x_val.shape)

    # Baseline model
    print("\n=== Training CIFAR-10 Baseline Model
(LR=0.001, Dropout=0.2, BS=32) ===")
    baseline_model =
build_model(learning_rate=0.001,
dropout_rate=0.2)
    baseline_history =
train_model(baseline_model, x_train, y_train,
x_val, y_val, batch_size=32, epochs=25)

    print("\nEvaluating baseline on test set...")
    baseline_loss, baseline_acc =
baseline_model.evaluate(x_test, y_test,
verbose=0)
    print(f"Baseline Test Loss: {baseline_loss:.4f},
Test Accuracy: {baseline_acc:.4f}")

    y_pred = baseline_model.predict(x_test,
verbose=0)
    y_pred_classes = np.argmax(y_pred, axis=1)
    y_true_classes = np.argmax(y_test, axis=1)
    baseline_report =
classification_report(y_true_classes,
y_pred_classes, target_names=classes,
output_dict=True, zero_division=0)

```

```

baseline_f1 = baseline_report['weighted
avg'] ['f1-score']

plot_results(baseline_history, title='CIFAR-10
Baseline EfficientNetB0')
save_confusion_matrix(y_true_classes,
y_pred_classes, classes,
'confusion_matrix_baseline_cifar10.png')

baseline_result = {
    'Model': 'Baseline',
    'Learning_Rate': 0.001,
    'Dropout_Rate': 0.2,
    'Batch_Size': 32,
    'Test_Loss': baseline_loss,
    'Test_Accuracy': baseline_acc,
    'Avg_Precision': baseline_report['weighted
avg'] ['precision'],
    'Avg_Recall': baseline_report['weighted
avg'] ['recall'],
    'Avg_F1_Score': baseline_f1
}

pd.DataFrame([baseline_result]).to_csv('cifar10_b
aseline_results.csv', index=False)
print("Baseline results saved to
cifar10_baseline_results.csv")

# Grid search
learning_rates = [0.001, 0.0005, 0.0001]
dropout_rates = [0.2, 0.3, 0.4]
batch_sizes = [32, 64, 128]

results = []
best_f1 = 0.0
best_params = None
best_model = None
best_y_pred = None
best_y_true = None

for lr, dr, bs in itertools.product(learning_rates,
dropout_rates, batch_sizes):
    print(f'\n=== Training CIFAR-10 with LR: {lr},
Dropout: {dr}, Batch Size: {bs} ===')

    try:
        model = build_model(learning_rate=lr,
dropout_rate=dr)

        history = train_model(model, x_train,
y_train, x_val, y_val, batch_size=bs, epochs=25)

```

```

        if (lr in [0.001, 0.0001] and dr in [0.2, 0.4]
and bs == 32) or (lr == 0.0005 and dr == 0.3 and
bs == 64):
            title = f'CIFAR-10
LR={lr}_DR={dr}_BS={bs}'
            plot_results(history, title)

            print("\nEvaluating on test set...")
            test_loss, test_acc =
model.evaluate(x_test, y_test, verbose=0)
            print(f"Test Loss: {test_loss:.4f}, Test
Accuracy: {test_acc:.4f}")

            y_pred = model.predict(x_test, verbose=0)
            y_pred_classes = np.argmax(y_pred,
axis=1)
            y_true_classes = np.argmax(y_test,
axis=1)
            report =
classification_report(y_true_classes,
y_pred_classes, target_names=classes,
output_dict=True, zero_division=0)
            f1_score = report['weighted
avg'] ['f1-score']

            if not results or f1_score > best_f1:
                best_f1 = f1_score
                best_params = (lr, dr, bs)
                best_model = model
                best_y_pred = y_pred_classes
                best_y_true = y_true_classes
                save_confusion_matrix(y_true_classes,
y_pred_classes, classes,
'confusion_matrix_best_cifar10.png')

            result = {
                'Learning_Rate': lr,
                'Dropout_Rate': dr,
                'Batch_Size': bs,
                'Test_Loss': test_loss,
                'Test_Accuracy': test_acc,
                'Avg_Precision': report['weighted
avg'] ['precision'],
                'Avg_Recall': report['weighted
avg'] ['recall'],
                'Avg_F1_Score': f1_score
            }
            results.append(result)

        except Exception as e:
            print(f"Error in evaluation for LR={lr},
DR={dr}, BS={bs}: {str(e)}")

```

```

print("\n=== Best Model Summary and Results
===")
if best_model and best_params:
    lr, dr, bs = best_params
    print(f"Best Model Parameters: LR={lr},
Dropout={dr}, Batch Size={bs}")
    best_model.summary()

    print("\nGenerating classification report for
best model...")
    report = classification_report(best_y_true,
best_y_pred, target_names=classes,
output_dict=True, zero_division=0)
    report_df =
pd.DataFrame(report).transpose()

report_df.to_csv('classification_report_best_cifar1
0.csv')
print("Classification report saved to
classification_report_best_cifar10.csv")

best_result = next(r for r in results if
r['Learning_Rate'] == lr and r['Dropout_Rate'] ==
dr and r['Batch_Size'] == bs)
print("\nBest Model Results:")
print(f"Learning Rate:
{best_result['Learning_Rate']}")
print(f"Dropout Rate:
{best_result['Dropout_Rate']}")
print(f"Batch Size:
{best_result['Batch_Size']}")
print(f"Test Loss:
{best_result['Test_Loss']:.4f}")
print(f"Test Accuracy:
{best_result['Test_Accuracy']:.4f}")
print(f"Average Precision:
{best_result['Avg_Precision']:.4f}")
print(f"Average Recall:
{best_result['Avg_Recall']:.4f}")
print(f"Average F1-Score:
{best_result['Avg_F1_Score']:.4f}")
else:
    print("No best model found.")

print("\nSaving CIFAR-10 grid search results to
CSV...")
results_df = pd.DataFrame(results)

results_df.to_csv('cifar10_grid_search_results.csv
', index=False)

```

```

print("Results saved to
cifar10_grid_search_results.csv")

if results_df.empty:
    print("Warning: CSV is empty! Check logs
and outputs.")
    print("Output files:",
os.listdir('/kaggle/working'))
else:
    print("\nFinal Grid Search Results for
CIFAR-10:")
    print(results_df.to_string(index=False))

# Main execution
if __name__ == '__main__':
    print("Starting CIFAR-10 EfficientNetB0
baseline and grid search with 8,000 test images
(800/class)...")
    run_cnn_grid_search()

```

2. STL -10 CNN Experimentation

```

import os
import warnings
import tensorflow as tf
import tensorflow_datasets as tfds # Added for
STL-10 loading
from tensorflow.keras import layers, models
from tensorflow.keras.preprocessing.image import
ImageDataGenerator
from tensorflow.keras.applications import
EfficientNetB0
from tensorflow.keras.applications.efficientnet
import preprocess_input
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import Callback
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import itertools
from sklearn.metrics import classification_report,
confusion_matrix
from sklearn.model_selection import
train_test_split
import time

# Suppress TensorFlow warnings (including
CUDA-related ones)
os.environ["TF_CPP_MIN_LOG_LEVEL"] = "3"
os.environ["TF_XLA_FLAGS"] =
"--tf_xla_auto_jit=-1"
tf.config.optimizer.set_jit(False)
warnings.filterwarnings("ignore",
category=UserWarning, module="keras")

# Set random seed for reproducibility
tf.random.set_seed(42)
np.random.seed(42)

# Define class names for STL-10
stl10_classes = ['airplane', 'bird', 'car', 'cat', 'deer',
'dog', 'horse', 'monkey', 'ship', 'truck']
NUM_CLASSES = 10

# Custom callback for epoch timing
class EpochTimeCallback(Callback):
    def on_epoch_begin(self, epoch, logs=None):
        self.start_time = time.time()

    def on_epoch_end(self, epoch, logs=None):
        end_time = time.time()
        duration = end_time - self.start_time

```

```

        print(f"Epoch {epoch + 1} took {duration:.2f}
seconds")

# Load and preprocess STL-10 using
tensorflow_datasets
def load_stl10_data():
    """
    Loads STL-10 with 5,000 train (500/class) and
    8,000 test (800/class) using tfds.
    """
    try:
        # Load STL-10 dataset using
tensorflow_datasets
        ds_train, ds_test = tfds.load('stl10',
split=['train', 'test'], as_supervised=True,
shuffle_files=True)

        # Convert to numpy arrays
        x_train, y_train = [], []
        for image, label in tfds.as_numpy(ds_train):
            x_train.append(image)
            y_train.append(label)
        x_train = np.array(x_train)
        y_train = np.array(y_train)

        x_test, y_test = [], []
        for image, label in tfds.as_numpy(ds_test):
            x_test.append(image)
            y_test.append(label)
        x_test = np.array(x_test)
        y_test = np.array(y_test)

        # Verify shapes
        assert x_train.shape[0] == 5000, f"Expected
5000 training samples, got {x_train.shape[0]}"
        assert x_test.shape[0] == 8000, f"Expected
8000 test samples, got {x_test.shape[0]}"

        # STL-10 labels are 1-10, convert to 0-9 for
consistency
        y_train = y_train - 1
        y_test = y_test - 1

        # Preprocess
        x_train = x_train.astype('float32')
        x_test = x_test.astype('float32')
        x_train = preprocess_input(x_train)
        x_test = preprocess_input(x_test)
        y_train = tf.keras.utils.to_categorical(y_train,
NUM_CLASSES)
        y_test = tf.keras.utils.to_categorical(y_test,
NUM_CLASSES)

```



```

# Verify class distribution in test set
class_counts =
np.bincount(np.argmax(y_test, axis=1))
print("STL-10 test set class distribution:",
class_counts)
assert all(count == 800 for count in
class_counts), "Uneven class distribution in
STL-10 test set"

print("STL-10 x_train shape:", x_train.shape,
"y_train shape:", y_train.shape)
print("STL-10 x_test shape:", x_test.shape,
"y_test shape:", y_test.shape)
return x_train, y_train, x_test, y_test
except Exception as e:
print(f"Error loading STL-10: {str(e)}")
return None, None, None, None

# Build EfficientNetB0 model
def build_model(learning_rate=0.001,
dropout_rate=0.2):
"""
Constructs EfficientNetB0 with
BatchNormalization.
"""

base_model =
EfficientNetB0(weights='imagenet',
include_top=False, input_shape=(224, 224, 3))
base_model.trainable = False

model = models.Sequential([
layers.Input(shape=(96, 96, 3)), # STL-10
images are 96x96
layers.Resizing(224, 224),
base_model,
layers.GlobalAveragePooling2D(),
layers.Dense(128, activation='relu',
kernel_initializer='he_normal'),
layers.BatchNormalization(),
layers.Dropout(dropout_rate),
layers.Dense(NUM_CLASSES,
activation='softmax')
])

model.compile(optimizer=Adam(learning_rate=lea
rning_rate, clipnorm=1.0, weight_decay=1e-4),
loss='categorical_crossentropy',
metrics=['accuracy'])
return model

```

```

# Data augmentation
def create_data_generator():
"""
Applies augmentation with contrast.
Rotation  $\pm 20^\circ$ , shifts 10%, flips, brightness [0.8,
1.2].
"""
return ImageDataGenerator(
rotation_range=20,
width_shift_range=0.1,
height_shift_range=0.1,
horizontal_flip=True,
brightness_range=[0.8, 1.2],
fill_mode='nearest'
)

# Train and evaluate model
def train_model(model, x_train, y_train, x_val,
y_val, batch_size=32, epochs=25):
"""
Trains with augmentation, early stopping, and
epoch timing.
"""
datagen = create_data_generator()
datagen.fit(x_train)

early_stopping =
tf.keras.callbacks.EarlyStopping(
monitor='val_loss', patience=5,
restore_best_weights=True
)
epoch_timer = EpochTimeCallback()

history = model.fit(
datagen.flow(x_train, y_train,
batch_size=batch_size),
validation_data=(x_val, y_val),
epochs=epochs,
callbacks=[early_stopping, epoch_timer],
verbose=1
)
return history

# Plot training results
def plot_results(history, title='STL-10
EfficientNetB0 Performance'):
"""
Generates loss/accuracy plots for report
visualization.
"""
plt.figure(figsize=(12, 4))
plt.subplot(1, 2, 1)

```

```
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'],
label='Validation Loss')
plt.title(f'{title} - Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
```

```
plt.subplot(1, 2, 2)
plt.plot(history.history['accuracy'], label='Train
Accuracy')
plt.plot(history.history['val_accuracy'],
label='Validation Accuracy')
plt.title(f'{title} - Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
```

```
plt.savefig(f'{title.lower().replace(" ", "_").png}')
plt.close()
```

Generate confusion matrix

```
def save_confusion_matrix(y_true, y_pred,
classes, filename):
```

```
    """
```

Saves confusion matrix for detailed class analysis.

```
    """
```

```
    cm = confusion_matrix(y_true, y_pred)
    plt.figure(figsize=(10, 8))
    plt.imshow(cm, interpolation='nearest',
cmap=plt.cm.Blues)
    plt.title('Confusion Matrix')
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)
    plt.xlabel('Predicted Label')
    plt.ylabel('True Label')
    plt.tight_layout()
    plt.savefig(filename)
    plt.close()
```

Run baseline and grid search for STL-10

```
def run_cnn_grid_search():
```

```
    """
```

Runs baseline and grid search for STL-10 with BatchNormalization.

Baseline: LR=0.001, Dropout=0.2, BS=32.

Grid: LR=[0.001, 0.0005, 0.0001],

Dropout=[0.2, 0.3, 0.4], BS=[32, 64, 128].

```
    """
```

```
x_train, y_train, x_test, y_test =
load_stl10_data()
classes = stl10_classes
```

```
if x_train is None:
    return
```

```
x_train, x_val, y_train, y_val = train_test_split(
    x_train, y_train, test_size=0.2,
random_state=42
)
print("STL-10 Train shape:", x_train.shape,
"Validation shape:", x_val.shape)
```

Baseline model

```
print("\n=== Training STL-10 Baseline Model
(LR=0.001, Dropout=0.2, BS=32) ===")
```

```
baseline_model =
build_model(learning_rate=0.001,
dropout_rate=0.2)
```

```
baseline_history =
train_model(baseline_model, x_train, y_train,
x_val, y_val, batch_size=32, epochs=25)
```

```
print("\nEvaluating baseline on test set...")
```

```
baseline_loss, baseline_acc =
baseline_model.evaluate(x_test, y_test,
verbose=0)
```

```
print(f"Baseline Test Loss: {baseline_loss:.4f},
Test Accuracy: {baseline_acc:.4f}")
```

```
y_pred = baseline_model.predict(x_test,
verbose=0)
```

```
y_pred_classes = np.argmax(y_pred, axis=1)
```

```
y_true_classes = np.argmax(y_test, axis=1)
```

```
baseline_report =
classification_report(y_true_classes,
y_pred_classes, target_names=classes,
output_dict=True, zero_division=0)
```

```
baseline_f1 = baseline_report['weighted
avg']['f1-score']
```

```
plot_results(baseline_history, title='STL-10
Baseline EfficientNetB0')
```

```
save_confusion_matrix(y_true_classes,
y_pred_classes, classes,
'confusion_matrix_baseline_stl10.png')
```

```
baseline_result = {
    'Model': 'Baseline',
    'Learning_Rate': 0.001,
    'Dropout_Rate': 0.2,
```

```

        'Batch_Size': 32,
        'Test_Loss': baseline_loss,
        'Test_Accuracy': baseline_acc,
        'Avg_Precision': baseline_report['weighted
avg']['precision'],
        'Avg_Recall': baseline_report['weighted
avg']['recall'],
        'Avg_F1_Score': baseline_f1
    }

pd.DataFrame([baseline_result]).to_csv('stl10_ba
seline_results.csv', index=False)
print("Baseline results saved to
stl10_baseline_results.csv")

# Grid search
learning_rates = [0.001, 0.0005, 0.0001]
dropout_rates = [0.2, 0.3, 0.4]
batch_sizes = [32, 64, 128]

results = []
best_f1 = 0.0
best_params = None
best_model = None
best_y_pred = None
best_y_true = None

for lr, dr, bs in itertools.product(learning_rates,
dropout_rates, batch_sizes):
    print(f'\n=== Training STL-10 with LR: {lr},
Dropout: {dr}, Batch Size: {bs} ===')

    try:
        model = build_model(learning_rate=lr,
dropout_rate=dr)

        history = train_model(model, x_train,
y_train, x_val, y_val, batch_size=bs, epochs=25)

        if (lr in [0.001, 0.0001] and dr in [0.2, 0.4]
and bs == 32) or (lr == 0.0005 and dr == 0.3 and
bs == 64):
            title = f'STL-10
LR={lr}_DR={dr}_BS={bs}'
            plot_results(history, title)

            print("\nEvaluating on test set...")
            test_loss, test_acc =
model.evaluate(x_test, y_test, verbose=0)
            print(f"Test Loss: {test_loss:.4f}, Test
Accuracy: {test_acc:.4f}")

```

```

y_pred = model.predict(x_test, verbose=0)
y_pred_classes = np.argmax(y_pred,
axis=1)
y_true_classes = np.argmax(y_test,
axis=1)

report =
classification_report(y_true_classes,
y_pred_classes, target_names=classes,
output_dict=True, zero_division=0)
f1_score = report['weighted
avg']['f1-score']

if not results or f1_score > best_f1:
    best_f1 = f1_score
    best_params = (lr, dr, bs)
    best_model = model
    best_y_pred = y_pred_classes
    best_y_true = y_true_classes
    save_confusion_matrix(y_true_classes,
y_pred_classes, classes,
'confusion_matrix_best_stl10.png')

result = {
    'Learning_Rate': lr,
    'Dropout_Rate': dr,
    'Batch_Size': bs,
    'Test_Loss': test_loss,
    'Test_Accuracy': test_acc,
    'Avg_Precision': report['weighted
avg']['precision'],
    'Avg_Recall': report['weighted
avg']['recall'],
    'Avg_F1_Score': f1_score
}
results.append(result)

except Exception as e:
    print(f"Error in evaluation for LR={lr},
DR={dr}, BS={bs}: {str(e)}")

print("\n=== Best Model Summary and Results
===")
if best_model and best_params:
    lr, dr, bs = best_params
    print(f"Best Model Parameters: LR={lr},
Dropout={dr}, Batch Size={bs}")
    best_model.summary()

    print("\nGenerating classification report for
best model...")

```

```

report = classification_report(best_y_true,
best_y_pred, target_names=classes,
output_dict=True, zero_division=0)
report_df =
pd.DataFrame(report).transpose()

report_df.to_csv('classification_report_best_stl10.
csv')
print("Classification report saved to
classification_report_best_stl10.csv")

best_result = next(r for r in results if
r['Learning_Rate'] == lr and r['Dropout_Rate'] ==
dr and r['Batch_Size'] == bs)
print("\nBest Model Results:")
print(f"Learning Rate:
{best_result['Learning_Rate']}")
print(f"Dropout Rate:
{best_result['Dropout_Rate']}")
print(f"Batch Size:
{best_result['Batch_Size']}")
print(f"Test Loss:
{best_result['Test_Loss']:.4f}")
print(f"Test Accuracy:
{best_result['Test_Accuracy']:.4f}")
print(f"Average Precision:
{best_result['Avg_Precision']:.4f}")
print(f"Average Recall:
{best_result['Avg_Recall']:.4f}")
print(f"Average F1-Score:
{best_result['Avg_F1_Score']:.4f}")
else:
print("No best model found.")

print("\nSaving STL-10 grid search results to
CSV...")
results_df = pd.DataFrame(results)

results_df.to_csv('stl10_grid_search_results.csv',
index=False)
print("Results saved to
stl10_grid_search_results.csv")

if results_df.empty:
print("Warning: CSV is empty! Check logs
and outputs.")
print("Output files:",
os.listdir('/kaggle/working'))
else:
print("\nFinal Grid Search Results for
STL-10:")
print(results_df.to_string(index=False))

```

```

# Main execution
if __name__ == '__main__':
print("Starting STL-10 EfficientNetB0 baseline
and grid search with 8,000 test images
(800/class)...")
run_cnn_grid_search()

```

3. CIFAR-10 (CNN FINE TUNING)

```

import tensorflow as tf
from tensorflow.keras import layers, models

```

```

from tensorflow.keras.applications import
EfficientNetB0
from tensorflow.keras.applications.efficientnet
import preprocess_input
from tensorflow.keras.preprocessing.image import
ImageDataGenerator
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import
ReduceLROnPlateau
import numpy as np
from sklearn.model_selection import
train_test_split

```

```

# Set random seed for reproducibility
tf.random.set_seed(42)
np.random.seed(42)

```

```

# Define number of classes
NUM_CLASSES = 10

```

```

# Load and preprocess CIFAR-10 with stratified
subsampling
def load_cifar10_data():
    (x_train, y_train), (x_test, y_test) =
tf.keras.datasets.cifar10.load_data()

```

```

    # Subsample to match original setup (5,000
train, 8,000 test)
    train_samples_per_class = 500
    x_train_sub, y_train_sub = [], []
    for cls in range(NUM_CLASSES):
        cls_indices = np.where(y_train[:, 0] == cls)[0]
        selected_indices =
np.random.choice(cls_indices,
train_samples_per_class, replace=False)

```

```

x_train_sub.append(x_train[selected_indices])

```

```

y_train_sub.append(y_train[selected_indices])
x_train = np.vstack(x_train_sub)
y_train = np.vstack(y_train_sub)

```

```

    test_samples_per_class = 800
    x_test_sub, y_test_sub = [], []
    for cls in range(NUM_CLASSES):
        cls_indices = np.where(y_test[:, 0] == cls)[0]
        selected_indices =
np.random.choice(cls_indices,
test_samples_per_class, replace=False)
        x_test_sub.append(x_test[selected_indices])
        y_test_sub.append(y_test[selected_indices])
    x_test = np.vstack(x_test_sub)

```

```

y_test = np.vstack(y_test_sub)

```

```

# Preprocess
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train = preprocess_input(x_train)
x_test = preprocess_input(x_test)
y_train = tf.keras.utils.to_categorical(y_train,
NUM_CLASSES)
y_test = tf.keras.utils.to_categorical(y_test,
NUM_CLASSES)

```

```

return x_train, y_train, x_test, y_test

```

```

# Data augmentation (adjusted for CIFAR-10)
def create_data_generator():

```

```

    return ImageDataGenerator(
        rotation_range=10,
        width_shift_range=0.1,
        height_shift_range=0.1,
        horizontal_flip=True,
        brightness_range=[0.8, 1.2],
        zoom_range=0.1,
        fill_mode='nearest'
    )

```

```

# Build and fine-tune model
def build_and_finetune_model():
    base_model =
EfficientNetB0(weights='imagenet',
include_top=False, input_shape=(224, 224, 3))

```

```

    # Freeze BatchNormalization layers
    for layer in base_model.layers:
        if isinstance(layer,
layers.BatchNormalization):
            layer.trainable = False

```

```

    # Initially freeze all layers, then unfreeze the
last 25 layers
    base_model.trainable = False
    for layer in base_model.layers[-25:]:
        layer.trainable = True

```

```

    # Build model with original architecture
    model = models.Sequential([
        layers.Input(shape=(32, 32, 3)), # CIFAR-10
native size
        layers.Resizing(224, 224),
        base_model,
        layers.GlobalAveragePooling2D(),

```



```

        layers.Dense(128, activation='relu',
kernel_initializer='he_normal'),
        layers.BatchNormalization(),
        layers.Dropout(0.2), # Best dropout rate
        layers.Dense(NUM_CLASSES,
activation='softmax')
    ])

    # Compile with a higher learning rate and
adjusted scheduler

model.compile(optimizer=Adam(learning_rate=1e
-3, clipnorm=1.0, weight_decay=1e-4),
        loss='categorical_crossentropy',
        metrics=['accuracy'])

    print("\nFine-tuning model for CIFAR-10 (final
attempt with stratified sampling)...")
    return model

# Train and evaluate model
def train_model(model, x_train, y_train, x_val,
y_val, x_test, y_test):
    datagen = create_data_generator()
    datagen.fit(x_train)

    early_stopping =
tf.keras.callbacks.EarlyStopping(
        monitor='val_loss', patience=5,
restore_best_weights=True
    )
    reduce_lr = ReduceLROnPlateau(
        monitor='val_loss', factor=0.3, patience=2,
min_lr=1e-6, verbose=1
    )

    model.fit(
        datagen.flow(x_train, y_train,
batch_size=128), # Best batch size
        validation_data=(x_val, y_val),
        epochs=40, # Increased epochs for
fine-tuning
        callbacks=[early_stopping, reduce_lr],
        verbose=1
    )

    loss, accuracy = model.evaluate(x_test, y_test,
verbose=0)
    print(f"CIFAR-10 Test Accuracy after final
fine-tuning with stratified sampling: {accuracy *
100:.2f}%")

```

```

# Generate classification report
y_pred = model.predict(x_test, verbose=0)
y_test_classes = np.argmax(y_test, axis=1)
y_pred_classes = np.argmax(y_pred, axis=1)

# Define class names for STL-10
class_names = ['airplane', 'bird', 'car', 'cat',
'deer', 'dog', 'horse', 'monkey', 'ship', 'truck']
print("\nClassification Report for CIFAR-10 Test
Set:")
print(classification_report(y_test_classes,
y_pred_classes, target_names=class_names))
return loss, accuracy

# Main fine-tuning script for CIFAR-10
def run_finetuning():
    # Load dataset
    print("Loading CIFAR-10...")
    x_train, y_train, x_test, y_test =
load_cifar10_data()
    x_train, x_val, y_train, y_val = train_test_split(
        x_train, y_train, test_size=0.2,
random_state=42, stratify=y_train
    )

    # Build and fine-tune model
    model = build_and_finetune_model()
    train_model(model, x_train, y_train, x_val,
y_val, x_test, y_test)

if __name__ == '__main__':
    run_finetuning()

```

4. STL-10 (CNN FINE TUNING)

```

import tensorflow as tf
import tensorflow_datasets as tfds
from tensorflow.keras import layers, models

```

```

from tensorflow.keras.applications import
EfficientNetB0
from tensorflow.keras.applications.efficientnet
import preprocess_input
from tensorflow.keras.preprocessing.image import
ImageDataGenerator
from tensorflow.keras.optimizers import Adam
import numpy as np
from sklearn.model_selection import
train_test_split
from sklearn.metrics import classification_report

# Set random seed for reproducibility
tf.random.set_seed(42)
np.random.seed(42)

# Define number of classes
NUM_CLASSES = 10

# Load and preprocess STL-10
def load_stl10_data():
    ds_train, ds_test = tfds.load('stl10', split=['train',
'test'], as_supervised=True, shuffle_files=True)
    x_train, y_train = [], []
    for image, label in tfds.as_numpy(ds_train):
        x_train.append(image)
        y_train.append(label - 1) # Adjust labels
from 1-10 to 0-9
    x_train, y_train = np.array(x_train),
np.array(y_train)

    x_test, y_test = [], []
    for image, label in tfds.as_numpy(ds_test):
        x_test.append(image)
        y_test.append(label - 1)
    x_test, y_test = np.array(x_test),
np.array(y_test)

    # Preprocess
    x_train = x_train.astype('float32')
    x_test = x_test.astype('float32')
    x_train = preprocess_input(x_train)
    x_test = preprocess_input(x_test)
    y_train = tf.keras.utils.to_categorical(y_train,
NUM_CLASSES)
    y_test = tf.keras.utils.to_categorical(y_test,
NUM_CLASSES)

    return x_train, y_train, x_test, y_test

# Data augmentation
def create_data_generator():

```

```

return ImageDataGenerator(
    rotation_range=20,
    width_shift_range=0.1,
    height_shift_range=0.1,
    horizontal_flip=True,
    brightness_range=[0.8, 1.2],
    fill_mode='nearest'
)

# Build and fine-tune model
def build_and_finetune_model():
    base_model =
EfficientNetB0(weights='imagenet',
include_top=False, input_shape=(224, 224, 3))

    # Freeze BatchNormalization layers
    for layer in base_model.layers:
        if isinstance(layer,
layers.BatchNormalization):
            layer.trainable = False

    # Initially freeze all layers, then unfreeze the
last 10 layers
    base_model.trainable = False
    for layer in base_model.layers[-10:]:
        layer.trainable = True

    # Build model with original architecture
    model = models.Sequential([
        layers.Input(shape=(96, 96, 3)), # STL-10
native size
        layers.Resizing(224, 224),
        base_model,
        layers.GlobalAveragePooling2D(),
        layers.Dense(128, activation='relu',
kernel_initializer='he_normal'),
        layers.BatchNormalization(),
        layers.Dropout(0.2), # Best dropout rate
        layers.Dense(NUM_CLASSES,
activation='softmax')
    ])

    # Compile with a higher learning rate for
fine-tuning
    model.compile(optimizer=Adam(learning_rate=1e
-4, clipnorm=1.0, weight_decay=1e-4),
        loss='categorical_crossentropy',
        metrics=['accuracy'])

    print("\nFine-tuning model for STL-10...")
    return model

```

```

# Train and evaluate model
def train_model(model, x_train, y_train, x_val,
y_val, x_test, y_test):
    datagen = create_data_generator()
    datagen.fit(x_train)

    early_stopping =
tf.keras.callbacks.EarlyStopping(
    monitor='val_loss', patience=5,
restore_best_weights=True
)

    model.fit(
        datagen.flow(x_train, y_train,
batch_size=128), # Best batch size
        validation_data=(x_val, y_val),
        epochs=40, # Increased epochs for
fine-tuning
        callbacks=[early_stopping],
        verbose=1
    )

    # Evaluate model
    loss, accuracy = model.evaluate(x_test, y_test,
verbose=0)
    print(f"STL-10 Test Accuracy after fine-tuning:
{accuracy * 100:.2f}%")

    # Generate classification report
    y_pred = model.predict(x_test, verbose=0)
    y_test_classes = np.argmax(y_test, axis=1)
    y_pred_classes = np.argmax(y_pred, axis=1)

    # Define class names for STL-10
    class_names = ['airplane', 'bird', 'car', 'cat',
'deer', 'dog', 'horse', 'monkey', 'ship', 'truck']
    print("\nClassification Report for STL-10 Test
Set:")
    print(classification_report(y_test_classes,
y_pred_classes, target_names=class_names))

    return loss, accuracy

# Main fine-tuning script for STL-10
def run_finetuning():
    # Load dataset
    print("Loading STL-10...")
    x_train, y_train, x_test, y_test =
load_stl10_data()
    x_train, x_val, y_train, y_val = train_test_split(

```

```

        x_train, y_train, test_size=0.2,
random_state=42
    )

    # Build and fine-tune model
    model = build_and_finetune_model()
    train_model(model, x_train, y_train, x_val,
y_val, x_test, y_test)

if __name__ == '__main__':
    run_finetuning()

```

5. CIFAR-10 CV Experimentation

```

import cv2
import numpy as np

```

```

import matplotlib.pyplot as plt
import os
from sklearn.cluster import MiniBatchKMeans
from sklearn.decomposition import PCA
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score,
classification_report
from sklearn.model_selection import
StratifiedShuffleSplit
from sklearn.preprocessing import
StandardScaler
import tensorflow_datasets as tfds

# Load and sample CIFAR-10 dataset with
stratified sampling
def load_data(dataset='cifar10'):
    if dataset == 'cifar10':
        ds, info = tfds.load('cifar10', split=['train',
'test'], as_supervised=True, with_info=True)
        train_ds, test_ds = ds[0], ds[1]

        train_images, train_labels = [], []
        test_images, test_labels = [], []
        for image, label in tfds.as_numpy(train_ds):
            train_images.append(image)
            train_labels.append(label)
        for image, label in tfds.as_numpy(test_ds):
            test_images.append(image)
            test_labels.append(label)
        train_images = np.array(train_images)
        train_labels = np.array(train_labels)
        test_images = np.array(test_images)
        test_labels = np.array(test_labels)

        # Sample 5000 train images from the 50,000
train set
        sss_train = StratifiedShuffleSplit(n_splits=1,
train_size=5000, random_state=42)
        train_indices, _ =
next(sss_train.split(train_images, train_labels))
        train_images = train_images[train_indices]
        train_labels = train_labels[train_indices]

        # Sample 8000 test images from the 10,000
test set
        sss_test = StratifiedShuffleSplit(n_splits=1,
train_size=8000, random_state=42)
        test_indices, _ =
next(sss_test.split(test_images, test_labels))
        test_images = test_images[test_indices]
        test_labels = test_labels[test_indices]

        expected_train_size, expected_test_size =
5000, 8000
        expected_per_class = [500, 800]

        else:
            raise ValueError("Dataset must be 'cifar10'.")

        train_images = train_images.astype(np.uint8)
        test_images = test_images.astype(np.uint8)

        print(f"Training set size: {len(train_images)}
images")
        print(f"Test set size: {len(test_images)}
images")
        train_class_counts = np.bincount(train_labels,
minlength=10)
        test_class_counts = np.bincount(test_labels,
minlength=10)
        print("Training class distribution:",
train_class_counts)
        print("Test class distribution:",
test_class_counts)

        if len(train_images) != expected_train_size or
len(test_images) != expected_test_size:
            raise ValueError(f"Unexpected dataset size.
Expected {expected_train_size} train and
{expected_test_size} test images.")
        if not all(count == expected_per_class[0] for
count in train_class_counts if count > 0) or not
all(count == expected_per_class[1] for count in
test_class_counts if count > 0):
            raise ValueError(f"Class distribution is not
balanced. Expected {expected_per_class[0]} per
class in train, {expected_per_class[1]} per class in
test.")

        return train_images, train_labels, test_images,
test_labels

# Extract SIFT features with scale-invariant
detection
def extract_sift_features(images,
contrast_threshold=0.04, edge_threshold=10):
    sift =
cv2.SIFT_create(contrastThreshold=contrast_thre
shold, edgeThreshold=edge_threshold)
    descriptors = []
    for img in images:
        gray = cv2.cvtColor(img,
cv2.COLOR_BGR2GRAY)

```

```

keypoints, des =
sift.detectAndCompute(gray, None)
    if des is not None and len(keypoints) > 0:
        descriptors.append(des)
    return np.vstack(descriptors) if descriptors else
np.array([])

```

```

# Compute PCA variance ratio with normalization
def compute_variance_ratio(images,
    contrast_threshold=0.04, edge_threshold=10,
    dataset='cifar10'):

```

```

    descriptors = extract_sift_features(images,
    contrast_threshold, edge_threshold)
    if len(descriptors) == 0:
        print("No descriptors found")
        return 0
    scaler = StandardScaler()
    descriptors = scaler.fit_transform(descriptors)
    pca = PCA(n_components=min(128,
    descriptors.shape[0]-1), random_state=42)
    pca.fit(descriptors)
    total_var = np.sum(pca.explained_variance_)
    if total_var == 0:
        print("Warning: Total variance is zero, check
data or normalization")
        return 0
    variance_ratio =
np.cumsum(pca.explained_variance_ratio_)
    plt.plot(range(1, len(variance_ratio) + 1),
    variance_ratio, 'b-')
    plt.title(f'Cumulative Explained Variance Ratio
({dataset.upper()})')
    plt.xlabel('Number of Components')
    plt.ylabel('Cumulative Explained Variance
Ratio')
    plt.grid(True)

```

```

plt.savefig(f'/content/variance_ratio_{dataset}_con
trast{contrast_threshold}_edge{edge_threshold}.p
ng')

```

```

plt.close()
    ninety_percent_idx = np.argmax(variance_ratio
    >= 0.9) + 1
    print(f'90% variance at ~{ninety_percent_idx}
components")
    return ninety_percent_idx

```

```

# Apply PCA and return fitted PCA object
def apply_pca(descriptors, n_components=59):
    pca = PCA(n_components=n_components,
    random_state=42)

```

```

    transformed_descriptors =
pca.fit_transform(descriptors)
    return transformed_descriptors, pca

```

```

# Build BoW vocabulary
def build_vocabulary(descriptors, k,
    batch_size=4096, use_pca=False, pca_dims=59):
    if use_pca:
        descriptors, pca = apply_pca(descriptors,
    n_components=pca_dims)
    else:
        pca = None
        kmeans = MiniBatchKMeans(n_clusters=k,
    n_init='auto', random_state=42,
    batch_size=batch_size)
        kmeans.fit(descriptors)
    return kmeans, pca

```

```

# Compute spatial histograms with SPM (L2
normalization only)
def compute_spatial_histograms(images,
    kmeans, pca=None, spm_levels=[1],
    contrast_threshold=0.04, edge_threshold=10):

```

```

    sift =
cv2.SIFT_create(contrastThreshold=contrast_thre
shold, edgeThreshold=edge_threshold)
    histograms = []
    for img in images:
        gray = cv2.cvtColor(img,
    cv2.COLOR_BGR2GRAY)
        keypoints, des =
sift.detectAndCompute(gray, None)
        if des is None or len(keypoints) == 0:

```

```

histograms.append(np.zeros(kmeans.n_clusters *
    sum(lvl**2 for lvl in spm_levels)))

```

```

        continue
    if pca is not None:
        des = pca.transform(des)
        words = kmeans.predict(des)

```

```

        h, w = gray.shape
        img_hist = []
        for grid_size in spm_levels:
            grid_h, grid_w = h // grid_size, w //
            grid_size
            level_hist = np.zeros((grid_size *
            grid_size, kmeans.n_clusters))
            for i in range(grid_size):
                for j in range(grid_size):
                    mask = []

```



```

        for kp_idx, kp in
enumerate(keypoints):
            kp_x, kp_y = kp.pt
            if (i * grid_h <= kp_y < (i + 1) *
grid_h) and (j * grid_w <= kp_x < (j + 1) * grid_w):
                mask.append(True)
            else:
                mask.append(False)
            mask = np.array(mask)
            grid_words = words[mask]
            if len(grid_words) > 0:
                hist, _ = np.histogram(grid_words,
bins=range(kmeans.n_clusters + 1),
density=True)
                level_hist[i * grid_size + j] = hist
            img_hist.append(level_hist.flatten())
            img_hist = np.concatenate(img_hist)

# L2 normalization
img_hist = img_hist /
(np.linalg.norm(img_hist) + 1e-10)
histograms.append(img_hist)

histograms = np.array(histograms)
return histograms

# Save results to a file
def save_results(results,
filename='/content/experiment_results_sift_cifar1
0.txt'):
    with open(filename, 'a') as f:
        for result in results:
            f.write(result + '\n')

# Save classification report to a file
def save_classification_report(report,
filename='/content/classification_report_cifar10.txt
'):
    with open(filename, 'w') as f:
        f.write(report)

# Main pipeline
def main():
    dataset = 'cifar10'
    print(f"\nProcessing {dataset.upper()}")
    train_images, train_labels, test_images,
test_labels = load_data(dataset=dataset)

    contrast_thresholds = [0.02, 0.04, 0.08]
    edge_thresholds = [7.5, 10, 12.5]

    best_accuracy = 0.0

```

```

    best_config = None
    best_predictions = None
    best_true_labels = None

    for contrast_threshold in contrast_thresholds:
        for edge_threshold in edge_thresholds:
            print(f"\nContrast Threshold:
{contrast_threshold}, Edge Threshold:
{edge_threshold}")
            pca_dims =
compute_variance_ratio(train_images,
contrast_threshold, edge_threshold,
dataset=dataset)
            if pca_dims == 0:
                continue

            vocab_sizes = [750, 1000, 1250]
            spm_configs = [[1]]

            results = []
            for k in vocab_sizes:
                print(f"\nVocabulary size: {k}")
                train_descriptors =
extract_sift_features(train_images,
contrast_threshold, edge_threshold)
                if len(train_descriptors) == 0:
                    print("No descriptors found")
                    continue
                kmeans, pca =
build_vocabulary(train_descriptors, k,
batch_size=4096, use_pca=True,
pca_dims=pca_dims)

                for spm in spm_configs:
                    print(f"\nSPM Levels: {spm},
Normalization: L2")
                    train_hist =
compute_spatial_histograms(
                        train_images, kmeans, pca,
spm_levels=spm,
contrast_threshold=contrast_threshold,
edge_threshold=edge_threshold
                    )
                    test_hist =
compute_spatial_histograms(
                        test_images, kmeans, pca,
spm_levels=spm,
contrast_threshold=contrast_threshold,
edge_threshold=edge_threshold
                    )

# SVM with rbf kernel only

```

```

for C in [0.1, 1.0, 10.0]:
    svm = SVC(C=C, kernel='rbf',
gamma='scale', random_state=42)
    svm.fit(train_hist, train_labels)
    pred = svm.predict(test_hist)
    acc = accuracy_score(test_labels,
pred)

    result = f"Dataset: {dataset},
Contrast: {contrast_threshold}, Edge:
{edge_threshold}, Vocab Size: {k}, SPM: {spm},
Norm: L2, SVM (Kernel=rbf, C={C},
Gamma=scale) Accuracy: {acc:.4f}"
    print(result)
    results.append(result)

    # Track best configuration
    if acc > best_accuracy:
        best_accuracy = acc
        best_config = result
        best_predictions = pred
        best_true_labels = test_labels

save_results(results)

# Generate and save classification report for
the best configuration
if best_config is not None:
    print(f"\nBest Configuration: {best_config}")
    print("\nClassification Report for Best
Configuration:")
    class_names = ['airplane', 'automobile', 'bird',
'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
    report =
classification_report(best_true_labels,
best_predictions, target_names=class_names)
    print(report)
    save_classification_report(report)

if __name__ == "__main__":
    main()

```

```

# Import libraries
import cv2
import numpy as np
import matplotlib.pyplot as plt
import os
from sklearn.cluster import MiniBatchKMeans
from sklearn.decomposition import PCA

```

```

from sklearn.svm import SVC
from sklearn.metrics import accuracy_score,
classification_report
from sklearn.preprocessing import
StandardScaler
from sklearn.model_selection import
StratifiedShuffleSplit
import tensorflow_datasets as tfds

# Load and subsample CIFAR-10 dataset with
stratified sampling
def load_data(dataset='cifar10'):
    if dataset != 'cifar10':
        raise ValueError("Dataset must be 'cifar10'.")

    # Load CIFAR-10 dataset
    ds, info = tfds.load('cifar10', split=['train', 'test'],
as_supervised=True, with_info=True)
    train_ds, test_ds = ds[0], ds[1]

    # Convert to numpy arrays
    train_images, train_labels = [], []
    for image, label in tfds.as_numpy(train_ds):
        train_images.append(image)
        train_labels.append(label)
    train_images = np.array(train_images)
    train_labels = np.array(train_labels)

    test_images, test_labels = [], []
    for image, label in tfds.as_numpy(test_ds):
        test_images.append(image)
        test_labels.append(label)
    test_images = np.array(test_images)
    test_labels = np.array(test_labels)

    # Subsample with stratification
    target_train_size, target_test_size = 5000,
8000
    samples_per_class_train,
samples_per_class_test = 500, 800

    # Stratified subsampling for training set
    sss_train = StratifiedShuffleSplit(n_splits=1,
train_size=target_train_size, random_state=42)
    train_idx, _ = next(sss_train.split(train_images,
train_labels))
    train_images = train_images[train_idx]
    train_labels = train_labels[train_idx]

    # Stratified subsampling for test set
    sss_test = StratifiedShuffleSplit(n_splits=1,
train_size=target_test_size, random_state=42)

```

```

test_idx, _ = next(sss_test.split(test_images,
test_labels))
test_images = test_images[test_idx]
test_labels = test_labels[test_idx]

# Convert to uint8 for OpenCV
train_images = train_images.astype(np.uint8)
test_images = test_images.astype(np.uint8)

# Verify class balance
print(f"Training set size: {len(train_images)}
images")
print(f"Test set size: {len(test_images)}
images")
train_class_counts = np.bincount(train_labels,
minlength=10)
test_class_counts = np.bincount(test_labels,
minlength=10)
print("Training class distribution:",
train_class_counts)
print("Test class distribution:",
test_class_counts)

if len(train_images) != target_train_size or
len(test_images) != target_test_size:
    raise ValueError(f"Unexpected dataset size.
Expected {target_train_size} train and
{target_test_size} test images.")
if not all(count == samples_per_class_train for
count in train_class_counts if count > 0):
    raise ValueError(f"Training class distribution
is not balanced. Expected
{samples_per_class_train} per class.")
if not all(count == samples_per_class_test for
count in test_class_counts if count > 0):
    raise ValueError(f"Test class distribution is
not balanced. Expected {samples_per_class_test}
per class.")

return train_images, train_labels, test_images,
test_labels

# Extract Harris keypoints and refine with LoG for
scale
def extract_harris_log_features(images,
block_size=3, k=0.04, sigma_values=[1.0, 1.5,
2.0], contrast_threshold=0.04,
edge_threshold=10):
    descriptors = []
    for img in images:
        gray = cv2.cvtColor(img,
cv2.COLOR_BGR2GRAY)

```

```

        # Harris corner detection
        harris = cv2.cornerHarris(gray,
blockSize=block_size, ksize=3, k=k)
        # Threshold Harris response to get initial
keypoints
        coords = np.where(harris > 0.01 *
harris.max())
        keypoints = [cv2.KeyPoint(float(x), float(y),
1.0) for y, x in zip(coords[0], coords[1])]

        if not keypoints:
            descriptors.append(np.array([]))
            continue

        # Scale selection with LoG
        best_sigma = None
        best_response = -float('inf')
        for sigma in sigma_values:
            log = cv2.GaussianBlur(gray, (0, 0),
sigma)
            log_response = np.zeros_like(gray,
dtype=float)
            for y, x in zip(coords[0], coords[1]):
                log_response[y, x] = log[y, x] if 0 <= y <
gray.shape[0] and 0 <= x < gray.shape[1] else 0
            max_response =
np.max(log_response[coords])
            if max_response > best_response:
                best_response = max_response
                best_sigma = sigma

        # Compute SIFT descriptors with refined
scale
        sift =
cv2.SIFT_create(contrastThreshold=contrast_thre
shold, edgeThreshold=edge_threshold)
        for kp in keypoints:
            kp.size = best_sigma * 2
            _, des = sift.compute(gray, keypoints)
            if des is not None and len(des) > 0:
                descriptors.append(des)
        return np.vstack(descriptors) if descriptors and
descriptors[0].size else np.array([])

# Compute PCA variance ratio with normalization
def compute_variance_ratio(images,
block_size=3, k=0.04, sigma_values=[1.0, 1.5,
2.0], contrast_threshold=0.04,
edge_threshold=10, dataset='cifar10'):
    descriptors =
extract_harris_log_features(images, block_size,

```

```

k, sigma_values, contrast_threshold,
edge_threshold)
    if len(descriptors) == 0:
        print("No descriptors found")
        return 0
    scaler = StandardScaler()
    descriptors = scaler.fit_transform(descriptors)
    pca = PCA(n_components=min(128,
descriptors.shape[0]-1), random_state=42)
    pca.fit(descriptors)
    total_var = np.sum(pca.explained_variance_)
    if total_var == 0:
        print("Warning: Total variance is zero, check
data or normalization")
        return 0
    variance_ratio =
np.cumsum(pca.explained_variance_ratio_)
    plt.plot(range(1, len(variance_ratio) + 1),
variance_ratio, 'b-')
    plt.title(f'Cumulative Explained Variance Ratio
({dataset.upper()})')
    plt.xlabel('Number of Components')
    plt.ylabel('Cumulative Explained Variance
Ratio')
    plt.grid(True)

plt.savefig(f'/kaggle/working/variance_ratio_{datas
et}_block{block_size}_k{k}_contrast{contrast_thre
shold}_edge{edge_threshold}.png')
    plt.close()
    ninety_percent_idx = np.argmax(variance_ratio
>= 0.9) + 1
    print(f'90% variance at ~{ninety_percent_idx}
components')
    return ninety_percent_idx

# Apply PCA and return fitted PCA object
def apply_pca(descriptors, n_components=59):
    pca = PCA(n_components=n_components,
random_state=42)
    transformed_descriptors =
pca.fit_transform(descriptors)
    return transformed_descriptors, pca

# Build BoW vocabulary
def build_vocabulary(descriptors, k,
batch_size=4096, use_pca=False, pca_dims=59):
    if use_pca:
        descriptors, pca = apply_pca(descriptors,
n_components=pca_dims)
    else:
        pca = None

```

```

kmeans = MiniBatchKMeans(n_clusters=k,
n_init='auto', random_state=42,
batch_size=batch_size)
kmeans.fit(descriptors)
return kmeans, pca

# Compute spatial histograms with SPM (L2
normalization only)
def compute_spatial_histograms(images,
kmeans, pca=None, spm_levels=[1],
block_size=3, k=0.04, sigma_values=[1.0, 1.5,
2.0], contrast_threshold=0.04,
edge_threshold=10):
    histograms = []
    for img in images:
        gray = cv2.cvtColor(img,
cv2.COLOR_BGR2GRAY)
        # Harris corner detection for keypoints
        harris = cv2.cornerHarris(gray,
blockSize=block_size, ksize=3, k=k)
        coords = np.where(harris > 0.01 *
harris.max())
        keypoints = [cv2.KeyPoint(float(x), float(y),
1.0) for y, x in zip(coords[0], coords[1])]

        if not keypoints:

            histograms.append(np.zeros(kmeans.n_clusters *
sum(lvl**2 for lvl in spm_levels)))
            continue

        # Scale selection with LoG
        best_sigma = None
        best_response = -float('inf')
        for sigma in sigma_values:
            log = cv2.GaussianBlur(gray, (0, 0),
sigma)
            log_response = np.zeros_like(gray,
dtype=float)
            for y, x in zip(coords[0], coords[1]):
                log_response[y, x] = log[y, x] if 0 <= y <
gray.shape[0] and 0 <= x < gray.shape[1] else 0
            max_response =
np.max(log_response[coords])
            if max_response > best_response:
                best_response = max_response
                best_sigma = sigma

        # Compute SIFT descriptors for this image
        sift =
cv2.SIFT_create(contrastThreshold=contrast_thre
shold, edgeThreshold=edge_threshold)

```

```

for kp in keypoints:
    kp.size = best_sigma * 2
_, des = sift.compute(gray, keypoints)

if des is None or len(des) == 0:

histograms.append(np.zeros(kmeans.n_clusters *
sum(lvl**2 for lvl in spm_levels)))
    continue

# Transform descriptors and predict words
for this image
    if pca is not None:
        des = pca.transform(des)
        words = kmeans.predict(des)

# Compute spatial histogram
h, w = gray.shape
img_hist = []
for grid_size in spm_levels:
    grid_h, grid_w = h // grid_size, w //
grid_size
    level_hist = np.zeros((grid_size *
grid_size, kmeans.n_clusters))
    for i in range(grid_size):
        for j in range(grid_size):
            mask = []
            for y, x in zip(coords[0], coords[1]):
                if (i * grid_h <= y < (i + 1) * grid_h)
and (j * grid_w <= x < (j + 1) * grid_w):
                    mask.append(True)
            else:
                mask.append(False)
            mask = np.array(mask)
            grid_words = words[mask] if
len(words) > 0 else np.array([])
            if len(grid_words) > 0:
                hist, _ = np.histogram(grid_words,
bins=range(kmeans.n_clusters + 1),
density=True)
                level_hist[i * grid_size + j] = hist
            img_hist.append(level_hist.flatten())
            img_hist = np.concatenate(img_hist)

# L2 normalization
img_hist = img_hist /
(np.linalg.norm(img_hist) + 1e-10)
histograms.append(img_hist)

return np.array(histograms)

```

Save results to a file

```

def save_results(results,
filename='/kaggle/working/experiment_results_ha
rris_sift_cifar10.txt'):
    with open(filename, 'a') as f:
        for result in results:
            f.write(result + '\n')

```

```

# Save classification report to a file
def save_classification_report(report,
filename='/kaggle/working/classification_report_h
arris_sift_cifar10.txt'):
    with open(filename, 'w') as f:
        f.write(report)

```

Main pipeline

```

def main():
    dataset = 'cifar10'
    print(f"\nProcessing {dataset.upper()}")
    train_images, train_labels, test_images,
test_labels = load_data(dataset=dataset)

```

```

    block_sizes = [2, 3]
    k_values = [0.04, 0.06]
    sigma_values_options = [[1.0, 1.5, 2.0]]
    contrast_thresholds = [0.02, 0.04, 0.08]
    edge_thresholds = [7.5, 10, 12.5]

```

```

    best_accuracy = 0.0
    best_config = None
    best_predictions = None
    best_true_labels = None

```

```

    for block_size in block_sizes:
        for k in k_values:
            for sigma_values in
sigma_values_options:
                for contrast_threshold in
contrast_thresholds:
                    for edge_threshold in
edge_thresholds:
                        print(f"\nBlock Size: {block_size},
k: {k}, Sigma Values: {sigma_values}, Contrast
Threshold: {contrast_threshold}, Edge Threshold:
{edge_threshold}")
                        pca_dims =
compute_variance_ratio(train_images,
block_size, k, sigma_values, contrast_threshold,
edge_threshold, dataset=dataset)
                        if pca_dims == 0:
                            continue

```

vocab_sizes = [750, 1000, 1250]

```

spm_configs = [[1]]

results = []
for k_vocab in vocab_sizes:
    print(f"\nVocabulary size:
{k_vocab}")
    train_descriptors =
extract_harris_log_features(train_images,
block_size, k, sigma_values, contrast_threshold,
edge_threshold)
    if len(train_descriptors) == 0:
        print("No descriptors found")
        continue
    kmeans, pca =
build_vocabulary(train_descriptors, k_vocab,
batch_size=4096, use_pca=True,
pca_dims=pca_dims)

    for spm in spm_configs:
        print(f"\nSPM Levels: {spm},
Normalization: L2")
        train_hist =
compute_spatial_histograms(
            train_images, kmeans,
pca, spm_levels=spm, block_size=block_size,
k=k,

sigma_values=sigma_values,
contrast_threshold=contrast_threshold,

edge_threshold=edge_threshold
        )
        test_hist =
compute_spatial_histograms(
            test_images, kmeans, pca,
spm_levels=spm, block_size=block_size, k=k,

sigma_values=sigma_values,
contrast_threshold=contrast_threshold,

edge_threshold=edge_threshold
        )

        for C in [0.1, 1.0, 10.0]:
            svm = SVC(C=C,
kernel='rbf', gamma='scale', random_state=42)
            svm.fit(train_hist,
train_labels)

            pred =
svm.predict(test_hist)
            acc =
accuracy_score(test_labels, pred)

```

```

result = f"Dataset:
{dataset}, Block Size: {block_size}, k: {k}, Sigma:
{sigma_values}, Contrast: {contrast_threshold},
Edge: {edge_threshold}, Vocab Size: {k_vocab},
SPM: {spm}, Norm: L2, SVM (Kernel=rbf, C={C},
Gamma=scale) Accuracy: {acc:.4f}"
print(result)
results.append(result)

if acc > best_accuracy:
    best_accuracy = acc
    best_config = result
    best_predictions = pred
    best_true_labels =

test_labels

save_results(results)

if best_config is not None:
    print(f"\nBest Configuration: {best_config}")
    print("\nClassification Report for Best
Configuration:")
    class_names = ['airplane', 'automobile', 'bird',
'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
    report =
classification_report(best_true_labels,
best_predictions, target_names=class_names)
    print(report)
    save_classification_report(report)

if __name__ == "__main__":
    main()

```

6. STL-10 CV Experimentation

```

# Import libraries
import cv2
import numpy as np
import matplotlib.pyplot as plt
import os

```

```

from sklearn.cluster import MiniBatchKMeans
from sklearn.decomposition import PCA
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score,
classification_report
from sklearn.preprocessing import
StandardScaler
import tensorflow_datasets as tfds
from sklearn.model_selection import
StratifiedShuffleSplit

# Load and verify STL-10 dataset
def load_data(dataset='stl10'):
    if dataset != 'stl10':
        raise ValueError("Dataset must be 'stl10'.")

    ds, info = tfds.load('stl10', split=['train', 'test'],
as_supervised=True, with_info=True)
    train_ds, test_ds = ds[0], ds[1]

    train_images, train_labels = [], []
    for image, label in tfds.as_numpy(train_ds):
        train_images.append(image)
        train_labels.append(label)
    train_images = np.array(train_images)
    train_labels = np.array(train_labels)

    test_images, test_labels = [], []
    for image, label in tfds.as_numpy(test_ds):
        test_images.append(image)
        test_labels.append(label)
    test_images = np.array(test_images)
    test_labels = np.array(test_labels)

    expected_train_size, expected_test_size =
5000, 8000
    expected_per_class = [500, 800]

    train_images = train_images.astype(np.uint8)
    test_images = test_images.astype(np.uint8)

    print(f"Training set size: {len(train_images)}
images")
    print(f"Test set size: {len(test_images)}
images")
    train_class_counts = np.bincount(train_labels,
minlength=10)
    test_class_counts = np.bincount(test_labels,
minlength=10)
    print("Training class distribution:",
train_class_counts)

```

```

print("Test class distribution:",
test_class_counts)

    if len(train_images) != expected_train_size or
len(test_images) != expected_test_size:
        raise ValueError(f"Unexpected dataset size.
Expected {expected_train_size} train and
{expected_test_size} test images.")
    if not all(count == expected_per_class[0] for
count in train_class_counts if count > 0) or not
all(count == expected_per_class[1] for count in
test_class_counts if count > 0):
        raise ValueError(f"Class distribution is not
balanced. Expected {expected_per_class[0]} per
class in train, {expected_per_class[1]} per class in
test.")

    return train_images, train_labels, test_images,
test_labels

# Extract Harris keypoints and refine with LoG for
scale
def extract_harris_log_features(images,
block_size=3, k=0.04, sigma_values=[1.0, 2.0,
4.0], contrast_threshold=0.04,
edge_threshold=10):
    descriptors = []
    for img in images:
        gray = cv2.cvtColor(img,
cv2.COLOR_BGR2GRAY)
        # Harris corner detection
        harris = cv2.cornerHarris(gray,
blockSize=block_size, ksize=3, k=k)
        # Threshold Harris response to get initial
keypoints
        coords = np.where(harris > 0.01 *
harris.max())
        keypoints = [cv2.KeyPoint(float(x), float(y),
1.0) for y, x in zip(coords[0], coords[1])]

        if not keypoints:
            descriptors.append(np.array([]))
            continue

    # Scale selection with LoG
    best_sigma = None
    best_response = -float('inf')
    for sigma in sigma_values:
        log = cv2.GaussianBlur(gray, (0, 0),
sigma)
        log_response = np.zeros_like(gray,
dtype=float)

```



```

        for y, x in zip(coords[0], coords[1]):
            log_response[y, x] = log[y, x] if 0 <= y <
gray.shape[0] and 0 <= x < gray.shape[1] else 0
            max_response =
np.max(log_response[coords])
            if max_response > best_response:
                best_response = max_response
                best_sigma = sigma

# Compute SIFT descriptors with refined
scale
sift =
cv2.SIFT_create(contrastThreshold=contrast_thre
shold, edgeThreshold=edge_threshold)
for kp in keypoints:
    kp.size = best_sigma * 2
    _, des = sift.compute(gray, keypoints)
    if des is not None and len(des) > 0:
        descriptors.append(des)
return np.vstack(descriptors) if descriptors and
descriptors[0].size else np.array([])

# Compute PCA variance ratio with normalization
def compute_variance_ratio(images,
block_size=3, k=0.04, sigma_values=[1.0, 2.0,
4.0], contrast_threshold=0.04,
edge_threshold=10, dataset='stl10'):
    descriptors =
extract_harris_log_features(images, block_size,
k, sigma_values, contrast_threshold,
edge_threshold)
    if len(descriptors) == 0:
        print("No descriptors found")
        return 0
    scaler = StandardScaler()
    descriptors = scaler.fit_transform(descriptors)
    pca = PCA(n_components=min(128,
descriptors.shape[0]-1), random_state=42)
    pca.fit(descriptors)
    total_var = np.sum(pca.explained_variance_)
    if total_var == 0:
        print("Warning: Total variance is zero, check
data or normalization")
        return 0
    variance_ratio =
np.cumsum(pca.explained_variance_ratio_)
    plt.plot(range(1, len(variance_ratio) + 1),
variance_ratio, 'b-')
    plt.title(f'Cumulative Explained Variance Ratio
({dataset.upper()})')
    plt.xlabel('Number of Components')

```

```

plt.ylabel('Cumulative Explained Variance
Ratio')
plt.grid(True)

plt.savefig(f'/content/Results/variance_ratio_{data
set}_block{block_size}_k{k}_contrast{contrast_thr
eshold}_edge{edge_threshold}.png')
plt.close()
    ninety_percent_idx = np.argmax(variance_ratio
>= 0.9) + 1
    print(f'90% variance at ~{ninety_percent_idx}
components")
    return ninety_percent_idx

# Apply PCA and return fitted PCA object
def apply_pca(descriptors, n_components=59):
    pca = PCA(n_components=n_components,
random_state=42)
    transformed_descriptors =
pca.fit_transform(descriptors)
    return transformed_descriptors, pca

# Build BoW vocabulary
def build_vocabulary(descriptors, k,
batch_size=4096, use_pca=False, pca_dims=59):
    if use_pca:
        descriptors, pca = apply_pca(descriptors,
n_components=pca_dims)
    else:
        pca = None
        kmeans = MiniBatchKMeans(n_clusters=k,
n_init='auto', random_state=42,
batch_size=batch_size)
        kmeans.fit(descriptors)
        return kmeans, pca

# Compute spatial histograms with SPM (L2
normalization only)
def compute_spatial_histograms(images,
kmeans, pca=None, spm_levels=[1],
block_size=3, k=0.04, sigma_values=[1.0, 2.0,
4.0], contrast_threshold=0.04,
edge_threshold=10):
    histograms = []
    for img in images:
        gray = cv2.cvtColor(img,
cv2.COLOR_BGR2GRAY)
        # Harris corner detection for keypoints
        harris = cv2.cornerHarris(gray,
blockSize=block_size, ksize=3, k=k)
        coords = np.where(harris > 0.01 *
harris.max())

```

```

keypoints = [cv2.KeyPoint(float(x), float(y),
1.0) for y, x in zip(coords[0], coords[1])]

if not keypoints:

histograms.append(np.zeros(kmeans.n_clusters *
sum(lvl**2 for lvl in spm_levels)))
    continue

# Scale selection with LoG
best_sigma = None
best_response = -float('inf')
for sigma in sigma_values:
    log = cv2.GaussianBlur(gray, (0, 0),
sigma)
    log_response = np.zeros_like(gray,
dtype=float)
    for y, x in zip(coords[0], coords[1]):
        log_response[y, x] = log[y, x] if 0 <= y <
gray.shape[0] and 0 <= x < gray.shape[1] else 0
    max_response =
np.max(log_response[coords])
    if max_response > best_response:
        best_response = max_response
        best_sigma = sigma

# Compute SIFT descriptors for this image
sift =
cv2.SIFT_create(contrastThreshold=contrast_thre
shold, edgeThreshold=edge_threshold)
    for kp in keypoints:
        kp.size = best_sigma * 2
        _, des = sift.compute(gray, keypoints)

    if des is None or len(des) == 0:

histograms.append(np.zeros(kmeans.n_clusters *
sum(lvl**2 for lvl in spm_levels)))
    continue

# Transform descriptors and predict words
for this image
    if pca is not None:
        des = pca.transform(des)
        words = kmeans.predict(des)

# Compute spatial histogram
h, w = gray.shape
img_hist = []
for grid_size in spm_levels:
    grid_h, grid_w = h // grid_size, w //
grid_size

```

```

        level_hist = np.zeros((grid_size *
grid_size, kmeans.n_clusters))
        for i in range(grid_size):
            for j in range(grid_size):
                mask = []
                for y, x in zip(coords[0], coords[1]):
                    if (i * grid_h <= y < (i + 1) * grid_h)
and (j * grid_w <= x < (j + 1) * grid_w):
                        mask.append(True)
                    else:
                        mask.append(False)
                mask = np.array(mask)
                grid_words = words[mask] if
len(words) > 0 else np.array([])
                if len(grid_words) > 0:
                    hist, _ = np.histogram(grid_words,
bins=range(kmeans.n_clusters + 1),
density=True)
                    level_hist[i * grid_size + j] = hist
                img_hist.append(level_hist.flatten())
            img_hist = np.concatenate(img_hist)

# L2 normalization
img_hist = img_hist /
(np.linalg.norm(img_hist) + 1e-10)
            histograms.append(img_hist)

    return np.array(histograms)

# Save results to a file
def save_results(results,
filename='/kaggle/working/experiment_results_ha
rris_sift_stl10.txt'):
    with open(filename, 'a') as f:
        for result in results:
            f.write(result + '\n')

# Save classification report to a file
def save_classification_report(report,
filename='/kaggle/working/classification_report_h
arris_sift_stl10.txt'):
    with open(filename, 'w') as f:
        f.write(report)

# Main pipeline
def main():
    dataset = 'stl10'
    print(f"\nProcessing {dataset.upper()}")
    train_images, train_labels, test_images,
test_labels = load_data(dataset=dataset)

    block_sizes = [3, 4]

```

```

k_values = [0.05, 0.07]
sigma_values_options = [[1.0, 2.0, 4.0]]
contrast_thresholds = [0.02, 0.04, 0.08]
edge_thresholds = [7.5, 10, 12.5]

best_accuracy = 0.0
best_config = None
best_predictions = None
best_true_labels = None

for block_size in block_sizes:
    for k in k_values:
        for sigma_values in
sigma_values_options:
            for contrast_threshold in
contrast_thresholds:
                for edge_threshold in
edge_thresholds:
                    print(f"\nBlock Size: {block_size},
k: {k}, Sigma Values: {sigma_values}, Contrast
Threshold: {contrast_threshold}, Edge Threshold:
{edge_threshold}")
                    pca_dims =
compute_variance_ratio(train_images,
block_size, k, sigma_values, contrast_threshold,
edge_threshold, dataset=dataset)
                    if pca_dims == 0:
                        continue

                    vocab_sizes = [750, 1000, 1250]
                    spm_configs = [[1]]

                    results = []
                    for k_vocab in vocab_sizes:
                        print(f"\nVocabulary size:
{k_vocab}")
                        train_descriptors =
extract_harris_log_features(train_images,
block_size, k, sigma_values, contrast_threshold,
edge_threshold)
                        if len(train_descriptors) == 0:
                            print("No descriptors found")
                            continue
                        kmeans, pca =
build_vocabulary(train_descriptors, k_vocab,
batch_size=4096, use_pca=True,
pca_dims=pca_dims)

                        for spm in spm_configs:
                            print(f"\nSPM Levels: {spm},
Normalization: L2")

```

```

train_hist =
compute_spatial_histograms(
    train_images, kmeans,
pca, spm_levels=spm, block_size=block_size,
k=k,

sigma_values=sigma_values,
contrast_threshold=contrast_threshold,

edge_threshold=edge_threshold
)
test_hist =
compute_spatial_histograms(
    test_images, kmeans, pca,
spm_levels=spm, block_size=block_size, k=k,

sigma_values=sigma_values,
contrast_threshold=contrast_threshold,

edge_threshold=edge_threshold
)

for C in [0.1, 1.0, 10.0]:
    svm = SVC(C=C,
kernel='rbf', gamma='scale', random_state=42)
    svm.fit(train_hist,
train_labels)
    pred =
svm.predict(test_hist)
    acc =
accuracy_score(test_labels, pred)
    result = f"Dataset:
{dataset}, Block Size: {block_size}, k: {k}, Sigma:
{sigma_values}, Contrast: {contrast_threshold},
Edge: {edge_threshold}, Vocab Size: {k_vocab},
SPM: {spm}, Norm: L2, SVM (Kernel=rbf, C={C},
Gamma=scale) Accuracy: {acc:.4f}"
    print(result)
    results.append(result)

    if acc > best_accuracy:
        best_accuracy = acc
        best_config = result
        best_predictions = pred
        best_true_labels =

test_labels

print(results)

if best_config is not None:
    print(f"\nBest Configuration: {best_config}")

```

```

print("\nClassification Report for Best
Configuration:")
class_names = ['airplane', 'automobile', 'bird',
'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
report =
classification_report(best_true_labels,
best_predictions, target_names=class_names)
print(report)
save_classification_report(report)

if __name__ == "__main__":
    main()

import cv2
import numpy as np
import matplotlib.pyplot as plt
import os
from sklearn.cluster import MiniBatchKMeans
from sklearn.decomposition import PCA
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score,
classification_report
from sklearn.model_selection import
train_test_split
from sklearn.preprocessing import
StandardScaler
import tensorflow_datasets as tfds
from sklearn.model_selection import
StratifiedShuffleSplit

# Load and sample STL-10 dataset with stratified
sampling
def load_data(dataset='stl10'):
    if dataset == 'stl10':
        ds, info = tfds.load('stl10', split=['train', 'test'],
as_supervised=True, with_info=True)
        train_ds, test_ds = ds[0], ds[1]

        train_images, train_labels = [], []
        for image, label in tfds.as_numpy(train_ds):
            train_images.append(image)
            train_labels.append(label)
        train_images = np.array(train_images)
        train_labels = np.array(train_labels)

        test_images, test_labels = [], []
        for image, label in tfds.as_numpy(test_ds):
            test_images.append(image)
            test_labels.append(label)
        test_images = np.array(test_images)
        test_labels = np.array(test_labels)

```

```

expected_train_size, expected_test_size =
5000, 8000
expected_per_class = [500, 800]

else:
    raise ValueError("Dataset must be 'stl10'.")

train_images = train_images.astype(np.uint8)
test_images = test_images.astype(np.uint8)

print(f"Training set size: {len(train_images)}
images")
print(f"Test set size: {len(test_images)}
images")
train_class_counts = np.bincount(train_labels,
minlength=10)
test_class_counts = np.bincount(test_labels,
minlength=10)
print("Training class distribution:",
train_class_counts)
print("Test class distribution:",
test_class_counts)

if len(train_images) != expected_train_size or
len(test_images) != expected_test_size:
    raise ValueError(f"Unexpected dataset size.
Expected {expected_train_size} train and
{expected_test_size} test images.")
if not all(count == expected_per_class[0] for
count in train_class_counts if count > 0) or not
all(count == expected_per_class[1] for count in
test_class_counts if count > 0):
    raise ValueError(f"Class distribution is not
balanced. Expected {expected_per_class[0]} per
class in train, {expected_per_class[1]} per class in
test.")

return train_images, train_labels, test_images,
test_labels

# Extract SIFT features with scale-invariant
detection
def extract_sift_features(images,
contrast_threshold=0.04, edge_threshold=10):
    sift =
cv2.SIFT_create(contrastThreshold=contrast_thre
shold, edgeThreshold=edge_threshold)
    descriptors = []
    for img in images:
        gray = cv2.cvtColor(img,
cv2.COLOR_BGR2GRAY)

```

```

keypoints, des =
sift.detectAndCompute(gray, None)
    if des is not None and len(keypoints) > 0:
        descriptors.append(des)
    return np.vstack(descriptors) if descriptors else
np.array([])

```

```

# Compute PCA variance ratio with normalization
def compute_variance_ratio(images,
contrast_threshold=0.04, edge_threshold=10,
dataset='stl10'):

```

```

    descriptors = extract_sift_features(images,
contrast_threshold, edge_threshold)
    if len(descriptors) == 0:
        print("No descriptors found")
        return 0
    scaler = StandardScaler()
    descriptors = scaler.fit_transform(descriptors)
    pca = PCA(n_components=min(128,
descriptors.shape[0]-1), random_state=42)
    pca.fit(descriptors)
    total_var = np.sum(pca.explained_variance_)
    if total_var == 0:
        print("Warning: Total variance is zero, check
data or normalization")
        return 0
    variance_ratio =
np.cumsum(pca.explained_variance_ratio_)
    plt.plot(range(1, len(variance_ratio) + 1),
variance_ratio, 'b-')
    plt.title(f'Cumulative Explained Variance Ratio
({dataset.upper()}')')
    plt.xlabel('Number of Components')
    plt.ylabel('Cumulative Explained Variance
Ratio')
    plt.grid(True)

```

```

plt.savefig(f'/content/variance_ratio_{dataset}_con
trast{contrast_threshold}_edge{edge_threshold}.p
ng')

```

```

    plt.close()
    ninety_percent_idx = np.argmax(variance_ratio
>= 0.9) + 1
    print(f'90% variance at ~{ninety_percent_idx}
components")
    return ninety_percent_idx

```

```

# Apply PCA and return fitted PCA object
def apply_pca(descriptors, n_components=59):
    pca = PCA(n_components=n_components,
random_state=42)

```

```

    transformed_descriptors =
pca.fit_transform(descriptors)
    return transformed_descriptors, pca

```

```

# Build BoW vocabulary
def build_vocabulary(descriptors, k,
batch_size=4096, use_pca=False, pca_dims=59):
    if use_pca:
        descriptors, pca = apply_pca(descriptors,
n_components=pca_dims)
    else:
        pca = None
        kmeans = MiniBatchKMeans(n_clusters=k,
n_init='auto', random_state=42,
batch_size=batch_size)
        kmeans.fit(descriptors)
    return kmeans, pca

```

```

# Compute spatial histograms with SPM (L2
normalization only)
def compute_spatial_histograms(images,
kmeans, pca=None, spm_levels=[1],
contrast_threshold=0.04, edge_threshold=10):

```

```

    sift =
cv2.SIFT_create(contrastThreshold=contrast_thre
shold, edgeThreshold=edge_threshold)
    histograms = []
    for img in images:
        gray = cv2.cvtColor(img,
cv2.COLOR_BGR2GRAY)
        keypoints, des =
sift.detectAndCompute(gray, None)
        if des is None or len(keypoints) == 0:

```

```

histograms.append(np.zeros(kmeans.n_clusters *
sum(lvl**2 for lvl in spm_levels)))

```

```

        continue
    if pca is not None:
        des = pca.transform(des)
        words = kmeans.predict(des)

```

```

        h, w = gray.shape
        img_hist = []
        for grid_size in spm_levels:
            grid_h, grid_w = h // grid_size, w //
grid_size
            level_hist = np.zeros((grid_size *
grid_size, kmeans.n_clusters))
            for i in range(grid_size):
                for j in range(grid_size):
                    mask = []

```

```

        for kp_idx, kp in
enumerate(keypoints):
            kp_x, kp_y = kp.pt
            if (i * grid_h <= kp_y < (i + 1) *
grid_h) and (j * grid_w <= kp_x < (j + 1) * grid_w):
                mask.append(True)
            else:
                mask.append(False)
            mask = np.array(mask)
            grid_words = words[mask]
            if len(grid_words) > 0:
                hist, _ = np.histogram(grid_words,
bins=range(kmeans.n_clusters + 1),
density=True)
                level_hist[i * grid_size + j] = hist
            img_hist.append(level_hist.flatten())
            img_hist = np.concatenate(img_hist)

# L2 normalization
img_hist = img_hist /
(np.linalg.norm(img_hist) + 1e-10)
histograms.append(img_hist)

histograms = np.array(histograms)
return histograms

# Save results to a file
def save_results(results,
filename='/content/experiment_results_sift_stl10.t
xt'):
    with open(filename, 'a') as f:
        for result in results:
            f.write(result + '\n')

# Save classification report to a file
def save_classification_report(report,
filename='/content/classification_report_stl10.txt'):
    with open(filename, 'w') as f:
        f.write(report)

# Main pipeline
def main():
    dataset = 'stl10'
    print(f"\nProcessing {dataset.upper()}")
    train_images, train_labels, test_images,
test_labels = load_data(dataset=dataset)

    contrast_thresholds = [0.02, 0.04, 0.08]
    edge_thresholds = [7.5, 10, 12.5]

    best_accuracy = 0.0
    best_config = None

```

```

    best_predictions = None
    best_true_labels = None

    for contrast_threshold in contrast_thresholds:
        for edge_threshold in edge_thresholds:
            print(f"\nContrast Threshold:
{contrast_threshold}, Edge Threshold:
{edge_threshold}")
            pca_dims =
compute_variance_ratio(train_images,
contrast_threshold, edge_threshold,
dataset=dataset)
            if pca_dims == 0:
                continue

            vocab_sizes = [750, 1000, 1250]
            spm_configs = [[1]]

            results = []
            for k in vocab_sizes:
                print(f"\nVocabulary size: {k}")
                train_descriptors =
extract_sift_features(train_images,
contrast_threshold, edge_threshold)
                if len(train_descriptors) == 0:
                    print("No descriptors found")
                    continue
                kmeans, pca =
build_vocabulary(train_descriptors, k,
batch_size=4096, use_pca=True,
pca_dims=pca_dims)

                for spm in spm_configs:
                    print(f"\nSPM Levels: {spm},
Normalization: L2")
                    train_hist =
compute_spatial_histograms(
                        train_images, kmeans, pca,
                        spm_levels=spm,
                        contrast_threshold=contrast_threshold,
                        edge_threshold=edge_threshold
                    )
                    test_hist =
compute_spatial_histograms(
                        test_images, kmeans, pca,
                        spm_levels=spm,
                        contrast_threshold=contrast_threshold,
                        edge_threshold=edge_threshold
                    )

            # SVM with rbf kernel only
            for C in [0.1, 1.0, 10.0]:

```

```

        svm = SVC(C=C, kernel='rbf',
gamma='scale', random_state=42)
        svm.fit(train_hist, train_labels)
        pred = svm.predict(test_hist)
        acc = accuracy_score(test_labels,
pred)

        result = f"Dataset: {dataset},
Contrast: {contrast_threshold}, Edge:
{edge_threshold}, Vocab Size: {k}, SPM: {spm},
Norm: L2, SVM (Kernel=rbf, C={C},
Gamma=scale) Accuracy: {acc:.4f}"
        print(result)
        results.append(result)

        # Track best configuration
        if acc > best_accuracy:
            best_accuracy = acc
            best_config = result
            best_predictions = pred
            best_true_labels = test_labels

    save_results(results)

    # Generate and save classification report for
the best configuration
    if best_config is not None:
        print(f"\nBest Configuration: {best_config}")
        print("\nClassification Report for Best
Configuration:")
        class_names = ['airplane', 'automobile', 'bird',
'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
        report =
classification_report(best_true_labels,
best_predictions, target_names=class_names)
        print(report)
        save_classification_report(report)

if __name__ == "__main__":
    main()

```